

¹ **Fixed-Parameter Tractability of**
² **Learning Small Decision Trees**
³ **(full paper)**

⁴ **Anonymous author**

⁵ **Anonymous affiliation**

⁶ —— **Abstract** ——

⁷ We consider the NP-hard problem of finding a smallest decision tree which represents a given partially
⁸ defined Boolean formula. We establish fixed-parameter tractability of the problem with respect to
⁹ the NLC-width of the instance. We formulate a dynamic programming procedure which utilizes
¹⁰ the NLC-decomposition of the instance. For this to work, we establish a succinct representation
¹¹ of partial solutions, so that the space and time requirements of each dynamic programming step
¹² remain bounded in terms of the NLC-width.

¹³ **2012 ACM Subject Classification** Theory of computation → Design and analysis of algorithms →
¹⁴ Parameterized complexity and exact algorithms → Fixed parameter tractability

¹⁵ **Keywords and phrases** parameterized complexity, NLC-width, rank-width, decision trees, partially
¹⁶ defined Boolean formulas

17 **1 Introduction**

18 Decision trees have proved to be extremely useful tools for the describing, classifying,
 19 generalizing data [18, 22, 25]. In this paper, we consider decision trees for *classification*
 20 *instances (CIs)*, consisting of a finite set E of *examples* (also called *feature vectors*) over a
 21 finite set F of *features*. Each example $e \in E$ is a function $e : F \rightarrow \{0, 1\}$ which determines
 22 whether the feature f is true or false for e . Moreover, E is given as a partition $E^+ \uplus E^-$ into
 23 positive and negative examples. For instance, examples could represent medical patients and
 24 features diagnostic tests; a patient is positive or negative corresponding to whether they have
 25 been diagnosed with a certain disease or not. CIs are also called *partially* or *incompletely*
 26 *defined Boolean functions*, as we can consider the features as Boolean variables, and examples
 27 as truth assignments that evaluate to 0 (for positive examples) or 1 (for negative examples).
 28 CIs have been studied as a key concept for the logical analysis of data and in switching
 29 theory [4, 6, 5, 7, 8, 17, 20].

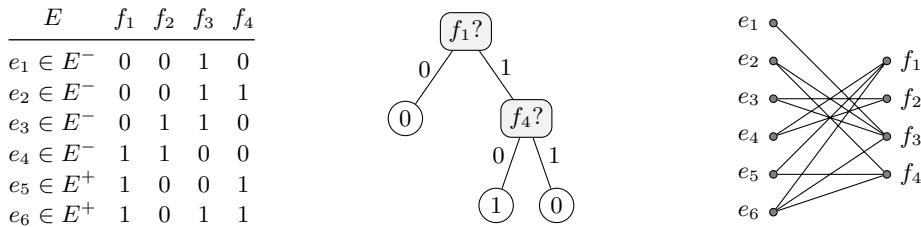
30 Because of their simplicity, decision trees are particularly attractive for providing in-
 31 terpretable models of the underlying CI, an aspect whose importance has been strongly
 32 emphasized over the recent years [10, 12, 15, 19, 21]. In this context, one prefers *small trees*,
 33 as they are easier to interpret and require fewer tests to make a classification. Small trees
 34 are also preferred in view of the parsimony principle (Occam's Razor) since small trees are
 35 expected to generalize better to new data [2]. However, finding a small decision tree, as
 36 formulated in the following decision problem, is NP-complete [16].

37 MINIMUM DECISION TREE SIZE (DTS): given a CI $E = E^+ \uplus E^-$ and an integer s ,
 38 is there a decision tree with at most s nodes for E ?

39 Given this complexity barrier, we propose a fixed-parameter algorithm for the problem,
 40 which exploits the input CI's hidden structure. The *incidence graph* of a CI is the bipartite
 41 graph $G_I(E)$ whose vertices are the examples on one side and the features on the other,
 42 where an example e is adjacent with a feature f if and only if $e(f) = 1$. Figure 1 shows a CI
 43 and a smallest decision tree for it, as well as the incidence graph.

44 Key to our algorithm are new notions for succinctly representing decision trees that
 45 correspond to subtrees of the incidence graph's tree decomposition. Based on that, we can
 46 carry out a dynamic programming (DP) procedure along the tree decomposition.

47 While the DP approach using treewidth is quite well understood and can often be quite
 48 easily designed for problems on graphs (or more generally problems whose solutions can be
 49 represented in terms of the graph for which the tree decomposition is given), the same DP
 50 approach can become rather involved if applied to problems whose solutions have no or only
 51 minor resemblance to the graph for which one is given a tree decomposition. Probably the
 52 most prominent example for this is the celebrated result by Bodlaender [3], where he uses a



■ **Figure 1** A CI $E = E^+ \uplus E^-$ with six examples and four features (left), a decision tree with 5 nodes that classifies E (middle), the incidence graph $G_I(E)$ (right).

53 DP approach on an approximate tree decomposition to compute the exact treewidth of a
 54 graph; here, the solutions are tree decompositions, which are complex structures that cannot
 55 easily be represented in terms of the graph. Other prominent examples include a DP approach
 56 to compute the exact treedepth [26] or clique-width [14] using an optimal tree decomposition.
 57 We face a similar problem, since solutions in our case are decision trees that do not bear
 58 any resemblance to the incidence graph for which we are given the tree decomposition. The
 59 main obstacle to overcome, therefore, is the design of the DP-records for our DP algorithm.
 60 That is, a record for a node b in a tree decomposition for the incidence graph of E needs
 61 to provide a compact representation of partial solutions, i.e. partial solutions in the sense
 62 that they represent the part of the solution for the whole instance E that corresponds to the
 63 sub-instance induced by all features and examples contained in the bags in the subtree of
 64 the tree decomposition rooted at the current node b . We overcome this obstacle in Section 3,
 65 where we also provide intuitive descriptions and motivation for the definition of the records
 66 (Subsection 3.1).

67 2 Preliminaries

68 2.1 Parameterized Complexity

69 We give some basic definitions of Parameterized Complexity and refer for a more in-depth
 70 treatment to other sources [9, 13]. Parameterized complexity considers problems in a two-
 71 dimensional setting, where a problem instance is a pair (I, k) , where I is the main part
 72 and k is the parameter. A parameterized problem is *fixed-parameter tractable* if there exists
 73 a computable function f such that instances (I, k) can be solved in time $f(k)\|I\|^{O(1)}$.

74 2.2 Graphs and NLC-width

75 We will assume that the reader is familiar with basic graph theory (see, e.g. [11, 1]). We
 76 consider (vertex and edge labelled) undirected graphs. Let $G = (V, E)$ be an undirected
 77 graph. We write $V(G) = V$ and $E(G) = E$ for the sets of vertices and edges of G , respectively.
 78 We denote an edge between $u \in V$ and $v \in V$ as $\{u, v\}$. For a set $V' \subseteq V$ of vertices we let
 79 $G[V']$ denote the graph induced by the vertices in V' , i.e. $G[V']$ has vertex set V' and edge
 80 set $E \cap \{\{u, v\} \mid u, v \in V'\}$ and we let $G - V'$ denote the graph $G[V \setminus V']$. For a set $E' \subseteq E$
 81 of edges we let denote $G - E'$ the graph with vertex set V and edge set $E \setminus E'$.

82 A *k-graph* is a pair (G, λ) , where $G = (V, E)$ is an undirected graph and $\lambda : V \rightarrow [k]$ is a
 83 *vertex label mapping* that labels every vertex $v \in V$ with a label $\lambda(v)$ from $[k]$. We call the
 84 *k-graph* consisting of exactly one vertex v (say, labeled by i) an *initial k-graph* and denote it
 85 by $i(v)$.

86 Node label control-width (*NLC-width*) is a graph parameter, defined as follows [28]: Let
 87 $k \in \mathbb{N}$ be a positive integer. An *k-NLC-expression tree* of a graph $G = (V, E)$ is a subcubic
 88 tree B , where every node b of B is associated with a *k-graph* (denoted by (G_b, λ_b)), such
 89 that:

- 90 1. Every leaf represents an initial *k-graph* $i(v)$ with $i \in [k]$ and $v \in V$.
- 91 2. Every non-leaf node b with one child c is a *relabeling node* and is associated with a
 92 relabelling function $R_b : [k] \rightarrow [k]$. Moreover, G_b is obtained from G_c after relabelling all
 93 vertices of G_c with label i to label $R_b(i)$ for every $i \in [k]$.
- 94 3. Every non-leaf node b with two children, i.e., a left child l and a right child r , is a *join*
 95 *node* and is associated with a *join matrix*, i.e., a binary $k \times k$ matrix M_b . Moreover,

(G_b, λ_b) is obtained from the disjoint union of (G_l, λ_l) and (G_r, λ_r) after adding an edge from all vertices labeled i in G_l to all vertices labeled j in G_r whenever $M_b[i, j] = 1$.

4. G is equal to the G_r for the root node r of B .

The NLC-width of a graph G , denoted by $nlcw(G)$, is the minimum k for which G has a k -NLC-expression tree. A k -NLC-expression tree is *nice* if every relabelling node has a relabelling function $R : [k] \rightarrow [k]$ such that for some $i, j \in [k]$, $R(i) = j$ and $R(\ell) = \ell$ for all $\ell \in [k] \setminus \{i\}$. Clearly, given a k -NLC-expression tree, a nice k -NLC-expression tree can be found in polynomial time; simply replace every relabelling node (that relabels more than one label at a time) by a sequence of relabelling nodes.

Let b be a node in a k -NLC-expression tree of a graph G . We denote by V_b the set of vertices of G_b . By the definition of a k -NLC-expression tree, if $u, v \in V_b$ have the same label in (G_b, λ_b) and $w \in V(G) \setminus V_b$, then u is adjacent to w in G if and only if v is.

Computing the NLC-width of a graph is NP-hard [?]. However, it is sufficient to use the algorithm of Seymour and Oum [?], which returns a c -expression for some $c \leq 2^{3cw(G)+2} - 1$ in $O(n^9 \log n)$ time, or the later improvements of Oum [24] and Hliněný and Oum [?] that provide cubic-time algorithms which yield a c -expression for some $c \leq 8^{cw(G)} - 1$ and $c \leq 2^{cw(G)+1} - 1$, respectively.

2.3 Classification Problems

An *example* e is a function $e : \text{feat}(e) \rightarrow \{0, 1\}$ defined on a finite set $\text{feat}(e)$ of *features*. For a set E of examples, we put $\text{feat}(E) = \bigcup_{e \in E} \text{feat}(e)$. We say that two examples e_1, e_2 *agree* on a feature f if $f \in \text{feat}(e_1)$, $f \in \text{feat}(e_2)$ and $e_1(f) = e_2(f)$. If $f \in \text{feat}(e_1)$, $f \in \text{feat}(e_2)$ but $e_1(f) \neq e_2(f)$, we say that the examples *disagree on* f .

A *classification instance* (CI) (also called a *partially defined Boolean function* [17]) $E = E^+ \uplus E^-$ is the disjoint union of two sets of examples, where for all $e_1, e_2 \in E$ we have $\text{feat}(e_1) = \text{feat}(e_2)$. The examples in E^+ are said to be *positive*; the examples in E^- are said to be *negative*. A set X of examples is *uniform* if $X \subseteq E^+$ or $X \subseteq E^-$; otherwise X is *non-uniform*.

Given a CI E , a subset $F \subseteq \text{feat}(E)$ is a *support set* of E if any two examples $e_1 \in E^+$ and $e_2 \in E^-$ disagree in at least one feature of F . Finding a smallest support set, denoted by $\text{MSS}(E)$, for a classification instance E is an NP-hard task [17, Theorem 12.2].

We define the *incidence graph* of E , denoted by $G_I(E)$, as the bipartite graph with partition $(E, \text{feat}(E))$ having an edge between an example $e \in E$ and a feature $f \in \text{feat}(e)$ if $f(e) = 1$.

2.4 Decision Trees

A *decision tree* (DT) (or *classification tree*) is a rooted tree T with vertex set $V(T)$ and arc set $A(T)$, where each non-leaf node (called a *test*) $v \in V(T)$ is labelled with a feature $\text{feat}(v)$, each non-leaf node v has exactly two out-going arcs, a *left arc* and a *right arc*, and each leaf is either a *positive* or a *negative* leaf. We write $\text{feat}(T) = \{v \in V(T) \mid \text{feat}(v)\}$.

Consider a CI E and a decision tree T with $\text{feat}(T) \subseteq \text{feat}(E)$. For each node v of T we define $E_T(v)$ as the set of all examples $e \in E$ such that for each left (right, respectively) arc (u, v) on the unique path from the root of T to v we have $e(\text{feat}(v)) = 0$ ($e(\text{feat}(v)) = 1$, respectively). T *correctly classifies* an example $e \in E$ if e is a positive (negative) example and $e \in E_T(v)$ for a positive (negative) leaf. We say that T *classifies* E (or simply that T is a DT for E) if T correctly classifies every example $e \in E$. See Figure 1 for an illustration of a CI, its incidence graph, and a DT that classifies E .

141 The size of T is its number of nodes, i.e. $|V(T)|$. We consider the following problem.

MINIMUM DECISION TREE SIZE (DTS)

142 Input: A classification instance E and an integer s .

Question: Is there a decision tree of size at most s for E ?

143 We now give some simple auxiliary lemmas that are required by our algorithm.

144 ▶ **Lemma 1.** Let A be a set of features of size a . Then the number of DTs of size at most s that use only features in A is at most a^{2s+1} and those can be enumerated in $\mathcal{O}(a^{2s+1})$ time.

145 **Proof.** We start by counting the number of trees T with n nodes that can potentially underlie a DT with n nodes. Note that there is one-to-one correspondence between trees T that underlie a DT with n nodes and unlabelled rooted ordered binary trees with n nodes (where ordered refers to an ordering of the at most 2 child nodes). Since it is known that the number of unlabelled rooted ordered binary trees with n nodes is equal to the n -th Catalan number C_n and that those trees can be enumerated in $\mathcal{O}(C_n)$ time [27], we already obtain that we can enumerate all of the at most C_n possible trees T underlying a DT of size n in $\mathcal{O}(C_n)$ time. Therefore, there are at most sC_s possible trees of size at most s that can underlie a DT with at most s nodes and those can be enumerated in $\mathcal{O}(sC_s)$ time. It now remains to bound the number of possible feature assignments $\text{feat}(f)$ for these trees as well as the number of possibilities for the leave nodes that can be either labelled positive or negative. Since we can assume that $a \geq 2$, we obtain that the number of possible feature assignments (and labellings of leaf-nodes) of a tree T with n nodes is at most a^n . Taking everything together, we obtain that there are at most $sC_s a^s \leq s4^s a^s \leq a^{2s+1}$ many DTs of size at most s using only features in A and those can be enumerated in $\mathcal{O}(a^{2s+1})$ time. ◀

161 ▶ **Lemma 2.** Let A be a set of features of size a . There are at most $a^{2^{a+1}+3}$ inclusion-wise minimal DTs using only features in A and these can be enumerated in $\mathcal{O}(a^{2^{a+1}+3})$ time.

162 **Proof.** Note that an inclusion-wise minimal DT T that uses only features in A has at most $2^a + 1$ nodes; this is because every feature appears at most once on every path T . Therefore, we obtain from Lemma 1 that the number of choices for T is at most $a^{2(2^a+1)+1} = a^{2^{a+1}+3}$. ◀

163 ▶ **Lemma 3.** Let E be a CI. Then one can decide whether E has a DT and if so output a DT of minimum size for E in time $\mathcal{O}((2^{|E|})^{4|E|-1})$.

164 **Proof.** Note first that $|\text{feat}(E)| \leq 2^{|E|}$ since we can assume that E does not contain two equivalent features. Moreover, E has a DT if and only if $\text{feat}(E)$ is a support set, which can be checked in time $\mathcal{O}(|E|^2 |\text{feat}(E)|)$ by checking, for every pair of positive and negative examples in E , whether there is a feature that distinguishes them. If this is not the case, we output **NO**, so assume that E has a DT. Note that any inclusion-wise minimal DT for E has at most $|E|$ leaves and therefore size at most $2|E| - 1$. We can therefore employ Lemma 1 to enumerate all inclusion-wise minimal potential DTs for E in time $\mathcal{O}((2^{|E|})^{2(2|E|-1)+1}) \in \mathcal{O}((2^{|E|})^{4|E|-1})$. For every such tree we then check whether it is indeed a DT for E and return a DT for E of minimum size found during this process. ◀

177 3 An FPT-Algorithm for NLC-width

178 In this section, we present our main result, i.e. we will show that DTS is fixed-parameter tractable parameterized by NLC-width.

180 ► **Theorem 4.** Let E be a CI, let B be an NLC-decomposition of width ω for $G_I(E)$, and
181 let s be an integer. Then, deciding whether E has a DT of size at most s is fixed-parameter
182 tractable parameterized by ω .

183 ► **Corollary 5.** DTS is fixed-parameter tractable parameterized by NLC-width.

todo: Due to
proposition ...

184 In principle, we will use a dynamic programming algorithm along the NLC-decomposition
185 (B, χ) of $G_I(E)$ that computes a set of records for every node b of B in a bottom-up manner.
186 Each record will represent an equivalence class of solutions (DTs) for the whole instance
187 restricted to the examples and features contained in the current subtree rooted in b , i.e.
188 the examples and features contained in $\chi(b)$. Before we continue with the formal notions
189 and definitions required to define the records, we want to illustrate the main ideas and
190 motivations. In what follows let B be an NLC-decomposition of $G_I(E)$ of width k . For
191 $b \in V(B)$, we write $\text{feat}(b)$ and $\text{exam}(b)$ for the sets $\chi(b) \cap \text{feat}(E)$ and $\chi(b) \cap E$, respectively.

192 3.1 Description of the Main Ideas Behind the Algorithm

193 Consider a node b of B . To simplify the presentation, we will sometime refer to the features
194 and examples in $\chi(B_b) \setminus \chi(b)$ as *forgotten* features and examples and we refer to the features
195 and examples in $(\text{feat}(E) \cup E) \setminus \chi(B_b)$ as *future* features and examples. We start with some
196 simple observations that follow immediately from the properties of tree decompositions.

- 197 ► **Observation 6.(1)** $e(f) = 0$ for every forgotten example $e \in \text{exam}(B_b) \setminus \text{exam}(b)$ and
198 future feature $f \in \text{feat}(E) \setminus \text{feat}(B_b)$,
199 (2) $e(f) = 0$ for every future example $e \in E \setminus \text{exam}(B_b)$ and forgotten feature $f \in \text{feat}(B_b) \setminus$
200 $\text{feat}(b)$;

todo: adjust to
NLC-width

201 **Proof.** Towards showing (1), let e be an example in $\text{exam}(B_b) \setminus \text{exam}(b)$ and let f be a
202 feature in $\text{feat}(E) \setminus \text{feat}(B_b)$. We claim that because (T, χ) is a tree decomposition of $G_I(E)$,
203 the graph $G_I(E)$ cannot contain an edge between e and f , which implies that $e(f) = 0$.
204 Suppose for a contradiction that this is not the case, i.e. $\{e, f\} \in E(G_I(E))$. Then, because
205 of property (T1) of a tree decomposition, there must exist a node b' such that $e, f \in \chi(b')$.
206 But then, if $b' \in V(B_b)$ we obtain that $f \notin \chi(b')$. Similarly, if $b' \in V(B \setminus B_b)$, we obtain
207 that $e \notin \chi(b')$ since otherwise e would violate property (T2) of a tree decomposition. This
208 completes the proof for (1); the proof for (2) is analogous. ◀

209 Informally, Observation 6 shows that forgotten examples cannot be distinguished by
210 future features and future examples cannot be distinguished by forgotten features. Consider
211 a DT T for E and a node b of B . For a set W containing features and examples from E , we
212 denote by $E[W]$ the sub-instance of E induced by the features and examples in W . Our aim
213 is to obtain a compact representation (represented by records) of the partial solution for the
214 sub-instance $E[\chi(B_b)]$ of E induced by the features and examples in $\chi(B_b)$ represented by T .

215 Intuitively, such a compact representation has to (1) represent a partial solution (DT)
216 for the examples in $\text{exam}(B_b)$ and (2) retain sufficient information about the structure of T
217 in order to decide whether it can be extended to a DT that also classifies the examples in
218 $E \setminus \text{exam}(B_b)$.

219 For illustration purposes let us first consider the simplified case that $\text{exam}(b) = \emptyset$. Because
220 of Observation 6 (1), this implies that every forgotten example goes to the left child of
221 any node t in T that is assigned a future feature. Therefore, under the assumption that
222 $\text{exam}(b) = \emptyset$ the DT T' obtained from T after:

- 223 ■ removing the subtree T_r of T for every right child r of a node t of T with $\text{feat}(t) \in$
224 $\text{feat}(E) \setminus \text{feat}(B_b)$ and replacing t with an edge from its parent in T to its left child in T

225 is a DT for $E[\chi(B_b)]$. Note that this means that under the rather strong assumption
226 that $\text{exam}(b) = \emptyset$, the part of T that takes care of the sub-instance $E[\chi(B_b)]$ is itself a DT
227 using only features in $\text{feat}(B_b)$; we will see later that unfortunately this is no longer the case
228 if $\text{exam}(b) \neq \emptyset$. Note that even though T' is a DT for $E[B_b]$, it does not yet constitute a
229 compact representation, since the number of features it uses in $\text{feat}(B_b) \setminus \text{feat}(b)$ is potentially
230 unbounded. However, we obtain from Observation 6 (2) that every future example will end
231 up in the left child of every node t of T' that is assigned a forgotten feature. This means
232 that to decide whether T' can be extended to a DT for the whole instance, the nodes that
233 are assigned forgotten features are not important. In fact, the only nodes in T' that can be
234 important for the classification of future examples are the nodes that are assigned features
235 in $\text{feat}(b)$. That is, it is sufficient to remember the DT T'' obtained from T' after:

- 236 ■ removing the subtree T_r of T' for every right child r of a node t of T' with $\text{feat}(t) \in$
237 $\text{feat}(B_b) \setminus \text{feat}(b)$ and replacing t with an edge from its parent in T' to its left child in T' .

238 Since the number of possible DT T'' is clearly bounded in terms of the number of features
239 in $\text{feat}(b)$ (and therefore in terms of the treewidth of $G_I(E)$), this would already give us the
240 compact representation that we are looking for. However, this only works in the case that
241 $\text{exam}(b) = \emptyset$, which is clearly not the case in general.

242 So let us now consider the general case with $\text{exam}(b) \neq \emptyset$. The first difference now is
243 that the part of T that takes care of the sub-instance $E[\chi(B_b)]$ is no longer a DT that only
244 uses features in $\text{feat}(B_b)$. In fact, it could even be the case that $E[\chi(B_b)]$ does not have a
245 DT, because there could exist examples in $\text{exam}(b)$ that can only be distinguished using
246 the features in $\text{feat}(E) \setminus \text{feat}(B_b)$. This means that we have to allow our partial solution for
247 $E[\chi(B_b)]$ to use future features. Fortunately, we do not need to know which exact future
248 feature is used by our partial solution but it suffices to know that a future feature is used and
249 how it behaves w.r.t. the examples in $\text{exam}(b)$; this is because Observation 6 (1) implies that
250 a future feature is used in a partial solution only for the purpose of distinguishing examples
251 in $\text{exam}(b)$. Moreover, because every forgotten example ends up in the left child of any node
252 t of T that uses a future feature, we only need to remember the left child for those nodes.
253 Also, we only need to remember occurrences of those nodes (using future features) if at least
254 one example in $\text{exam}(b)$ ends up in the right child of such a node; otherwise the node has
255 no influence on the classification of examples in $\text{exam}(B_b)$. Finally, we cannot simply forget
256 nodes that use forgotten features (as we could in the case that $\text{exam}(b) = \emptyset$). This is because
257 we need to know exactly where the examples in $\text{exam}(b)$ end up at. For instance, if such
258 an example in $\text{exam}(b)$ ends up in the right child of a node using a future feature, we need
259 to know that this is the case because this means that the example has to be classified in
260 this place at a later stage of the algorithm. Nevertheless, we do not need to remember all
261 occurrences of nodes using forgotten features, but only those for which there is at least one
262 example in $\text{exam}(b)$ that ends up in the right child of the node. Similarly, we do not need
263 to remember the exact forgotten feature that is used but only how it behaves towards the
264 examples in $\text{exam}(b)$. In summary, we only need to remember the full information about
265 the nodes of T that use a feature in $\text{feat}(b)$. For all other nodes, i.e. nodes that use either
266 forgotten or future features, we only need to remember such a node, if at least one example
267 in $\text{exam}(b)$ ends up in its right child. Moreover, even if this is the case, we only need to
268 remember the following for such nodes:

- 269 ■ whether it uses a future or a forgotten feature and

270 ■ how it behaves w.r.t. the examples in $\text{exam}(b)$.

271 With these ideas in mind, we are now ready to provide a formal definition of the compact
 272 representation of the part of T that takes care of the sub-instance $E[\chi(B_b)]$.

273 3.2 Formal Definition of Records and Preliminary Results

274 In the following, let E be a CI and let B be a k -NLC-expression tree for $G_I(E)$. Consider a
 275 node b of B . Recall that b is either a leaf node associated with a k -graph $i(v)$, a relabelling
 276 node with 1 child and with relabelling function R_b , or a join node with a left child, a right
 277 child and a join matrix M_b . Moreover, recall that (G_b, λ_b) is the k -graph associated with b
 278 (whose unlabeled version is a subgraph of G) and V_b is the set of vertices of G_b . Additionally,
 279 we will use the following notation. We denote by $\text{feat}(b)$ the set $V_b \cap \text{feat}(E)$ of features in
 280 V_b and by $\text{exam}(b)$ the set $V_b \cap E$ of examples in V_b .

281 Consider a node b of B . Let L be a set of labels (usually $L = [k]$). For a subset $L' \subseteq L$,
 282 we denote by $\overline{L'}$ the set $L \setminus L'$. For a label $l \in L$, we introduce a new feature f_l , which we
 283 will call a *forgotten feature*. Moreover, for a subset $L' \subseteq L$ of labels, we introduce a new
 284 feature $f_{L'}$, which we call an *future (or introduce) feature*. Let $F_L = \{f_l \mid l \in L\}$ be the set
 285 of all forgotten features and let $I_L = \{f_{L'} \mid L' \subseteq L\}$ be the set of all future features w.r.t. L .
 286 To distinguish features in $\text{feat}(E)$ from forgotton and future features, we will refer to them
 287 as *real features*.

288 Let T be a decision tree and $t \in V(T)$. We say that a node t_A is a *left/right ancestor*
 289 of t if t is contained in the subtree of T rooted at the left/right child of t_A . We denote by
 290 $\text{anc}_L(t)/\text{anc}_R(t)$ the set of all left/right ancestors of t in T . We denote by $\text{anc}(t)$ the set of
 291 all *ancestors* of t in T , i.e., $\text{anc}(t) = \text{anc}_L(t) \cup \text{anc}_R(t)$.

292 Let T be a decision tree and $t \in V(T)$ be an inner node of T with left child l , right child
 293 r , and parent p . We say that T' is obtained from T after *left/right-contracting* t if T' is the
 294 decision tree obtained from T after removing t together with all nodes in T_r/T_l and adding
 295 the edge between p and l/r ; if t has no parent then no edge is added.

296 We say that T is a *decision tree* for b , if T is a decision tree for $\text{exam}(b)$ that uses only
 297 the features in $\text{feat}(b)$. We say that an inner node $t \in V(T)$ is *left/right redundant* in T if
 298 $\text{feat}(t) \in \text{feat}(\text{anc}_L(t))/\text{feat}(t) \in \text{feat}(\text{anc}_R(t))$. We say that t is redundant if it is either left
 299 redundant or right redundant. Intuitively, a node t is left/right redundant if all examples
 300 that end up at t , i.e., the examples $E_T(t)$, go the left/right child of t in T . Therefore, if t
 301 is left/right redundant in T , then the tree obtained after left/right-contracting t is still a
 302 decision tree.

303 We say that T is a *decision tree template* for b if T is a decision tree for $\text{exam}(b)$ that can
 304 additionally use the future features in $I_{[k]}$. Here, we assume that a future feature $f_{L'} \in I_{[k]}$
 305 for some $L' \subseteq [k]$ is 1 at an example $e \in \text{exam}(b)$ if $\lambda_b(e) \in L'$ and otherwise it is 0. We say
 306 that a decision tree template is *complete* if it does not use any features in $I_{[k]}$, otherwise
 307 we say that it is *incomplete*. Informally, the role of the future features in a decision tree
 308 template is provide spaceholders for the features in $\text{feat}(E) \setminus \text{feat}(b)$. Because all of those
 309 features behave the same w.r.t. to examples in $\text{exam}(b)$ having the same label, they can
 310 be charactericed by the set of labels for which those features are 1. Let T be a decision
 311 tree template for b and let $t \in V(T)$. We denote by $A(t)$ the set of *filtered labels* for t , i.e.,
 312 $A(t) = (\bigcap_{f_{L'} \in \text{feat}(\text{anc}_L(t)) \cap I_{[k]}} \overline{L'}) \cap (\bigcap_{f_{L'} \in \text{feat}(\text{anc}_R(t)) \cap I_{[k]}} L')$. Informally, $A(t)$ is the set of all
 313 labels $l \in [k]$ such that an example e with label l would end up at t , if only the effect of
 314 the future features on the path to t is considered. We say that t with $f_{L'} = \text{feat}(t) \in I_{[k]}$ is
 315 *left/right redundant* in T if $A(t) \subseteq L'/A(t) \subseteq \overline{L'}$. We say that t is *redundant* if it is either

definition of new
features

316 left-redundant or right-redundant. Intuitively, t is left/right redundant if all examples that
 317 can reach t (considering the influence of the future features only) end up in the left/right
 318 child of t . This also implies that if t is left/right redundant then the decision tree obtained
 319 after left/right contracting t is equivalent with T (all examples end up in the same leaves).

320 We say that T is a *decision tree skeleton* for b if T is a decision tree that can only use
 321 features in $F_{[k]} \cup I_{[k]}$. Note that because of the features $F_{[k]}$, whose behaviour w.r.t. the
 322 examples in $\text{exam}(b)$ is not defined, the behaviour w.r.t. the examples in $\text{exam}(b)$ of such a
 323 DT skeleton is not necessarily defined. Nevertheless, the behaviour of a feature f_l in $F_{[k]}$ is
 324 well-defined w.r.t. to the examples in $\text{exam}(E) \setminus \text{exam}(b)$, i.e., it behaves the same as any
 325 feature in $\text{feat}(b)$ with label l . Intuitively, decision tree skeletons are obtained from decision
 326 tree templates after replacing every feature f in $\text{feat}(b)$ with its label $\lambda_b(f)$. This allows us to
 327 further compress the information contained in decision tree templates, while still keeping the
 328 information about how the decision tree template behaves w.r.t. future examples in $\text{exam}(b)$.
 329 In particular, decision tree skeletons will form the main information stored by our records.

330 Let T be a decision tree skeleton and $t \in V(T)$. Similarly as we did for decision tree
 331 templates, we say that T is *complete* if it uses no future features and otherwise we say that it
 332 is incomplete. We say that an inner node t with $f_l = \text{feat}(t) \in F_{[k]}$ is *left/right redundant* in
 333 T if $f_l \in \text{feat}(\text{anc}_L(t)) / f_l \in \text{feat}(\text{anc}_R(t))$. Similarly, as for decision tree (templates), if t is
 334 left/right redundant, then we can left/right contract t without changing the properties of T .

335 Let T be a decision tree (skeleton/template). Then, we denote by $r(T)$ the decision tree
 336 obtained from T after left/right contracting every left/right redundant node of T . Note that
 337 if T is a decision tree (skeleton/template) for b , then so is $r(T)$.

338 ▶ **Observation 7.** *Let T be a decision tree skeleton/template for b . Then, so is $r(T)$.*

a short proof

339 **Proof.**

340 We say that T is *reduced* if $r(T) = T$.

341 ▶ **Lemma 8.** *Let T be a reduced decision tree (skeleton/template) using at most a real
 342 features, b forgotten features, and c future features. Then, T has size at most ?.*

todo

343 **Proof.**

definition of
344 relabelling

345 Let T be a decision tree. A *feature relabelling* for T is a function $\alpha : F' \rightarrow \text{feat}(E) \cup$
 346 $F_L \cup I_L$, where $F' \subseteq \text{feat}(T)$ and L is some set of labels (usually $L = [k]$). With a
 347 slight abuse of notation, we denote by $\alpha(T)$, the decision tree obtained after relabeling all
 348 features in F' (used by T) according to α , i.e., $\alpha(T)$ is obtained from T after replacing the
 349 feature assignment function $\text{feat}_T(t)$ for T with the function $\text{feat}_{\alpha(T)}(t)$ defined by setting
 350 $\text{feat}_{\alpha(T)}(t) = \alpha(\text{feat}_T(t))$ if $\text{feat}(t) \in F'$ and $\text{feat}_{\alpha(T)}(t) = \text{feat}_T(t)$, otherwise. We say that
 351 two feature relabellings $\alpha_1 : F_1 \rightarrow \text{feat}(E) \cup F_L \cup I_L$ and $\alpha_2 : F_2 \rightarrow \text{feat}(E) \cup F_L \cup I_L$ are
 352 *compatible* if they agree on their shared domain $F_1 \cap F_2$.

353 We denote by α_b^s the *standard feature relabelling* for b , i.e., the function $\alpha_b^s : \text{feat}(b) \rightarrow [k]$
 354 defined by setting $\alpha_b^s(f) = \lambda_b(f)$ for every $f \in \text{feat}(b)$.

Semantics of
records

355 We are now ready to define the records and their semantics. A *record* for b is a pair (T, s)
 356 such that T is a reduced decision tree skeleton for b and s is a natural number. We say that a
 357 record (T, s) is *valid* for b if s is the minimum number such that there is a (reduced) decision
 358 tree template T' for b such that $r(\alpha_b^s(T')) = T$ and $s = |V(T') \setminus V(T)|$. We denote by $\mathcal{R}(b)$
 the set of all valid records for b . The following corollary follows immediately from Lemma 8.

359 ▶ **Corollary 9.** $|\mathcal{R}(b)| \leq ?$

360 Note that E has a DT of size at most s if and only if $\mathcal{R}(r)$ contains a record (T, s) such that
361 T is complete, where r is the root of B

362 ▶ **Lemma 10.** *Let T be a decision tree and let α be a feature relabelling for T . Then,
363 $r(\alpha(T)) = r(\alpha(r(T)))$.*

auxiliary
properties
of
feature
relabelings
and reductions

364 ▶ **Observation 11.** *Let T be a decision tree and let α_1 and α_2 be two compatible feature
365 relabelling for T . Then, $\alpha_1\alpha_2(T) = \alpha_2\alpha_1(T)$.*

366 3.3 Proof to the Main Result

367 We will now show that we can compute $\mathcal{R}(b)$ for every of the 3 node types of a nice k -NLC
368 expression tree provided that $\mathcal{R}(c)$ has already been computed for every child c of b .

369 ▶ **Lemma 12** (leaf node). *Let $b \in V(B)$ be a leaf node. Then $\mathcal{R}(b)$ can be computed in time
370 ??.*

371 **Proof.** Let $i(v)$ be the initial k -graph associated with b . If v is a feature, then $\mathcal{R}(b)$ contains
372 all records $(T, 0)$ such that T is a reduced decision tree skeleton for b using only the features
373 in $\{f_{\lambda(v)}\} \cup I_{[k]}$. The correctness in this case follows because V_b contains no examples and
374 therefore every reduced decision tree skeleton constitutes a valid record for b . Moreover, the
375 run-time follows from Lemma ??, since the time required to enumerate all those reduced
376 decision tree skeletons is at most $\mathcal{O}(?)$.

377 If, on the other hand v is an example, then $\mathcal{R}(b)$ contains all records $(T, 0)$ such that T
378 is a reduced decision tree skeleton for b using only the features in $I_{[k]}$ and which correctly
379 classify v . Because of Lemma ??, those can be enumerated in time $\mathcal{O}(?)$ and checking for
380 each of those whether it correctly classifies v can be achieved in time $\mathcal{O}(?)$.

◀
todo: show
correctness

382 ▶ **Lemma 13** (join node). *Let $b \in V(B)$ be a join node. Then $\mathcal{R}(b)$ can be computed in time
383 $\mathcal{O}(k(2k + 2^k + 2)2^{6k+1})$.*

384 **Proof.** Let b_L and b_R be the left and right child of b in B , respectively.

385 Let M_b be the join matrix for the node b , i.e., M_b is a $k \times k$ binary matrix. For every
386 label $i \in [k]$, let $A_{i,*} = \{j \in [k] \mid M_b[i, j] = 1\}$ and $A_{*,i} = \{j \in [k] \mid M_b[j, i] = 1\}$.

387 To distinguish between forgotten features from the left and the right subtree, we introduce
388 the left i_L and the right version i_R for every label $i \in [k]$. With a slight abuse of notation,
389 we also denote by $[k_L]$ be the set $\{1_L, \dots, k_L\}$ of (left) labels and we denote by $[k_R]$ be the
390 set $\{1_R, \dots, k_R\}$ of (right) labels.

391 To compute the set $\mathcal{R}(b)$ of valid record for b , we first enumerate all reduced DT skeletons
392 T using features in $[k_L] \cup [k_R] \cup I_{[k]}$. Because of Lemma 17, those can be enumerated in time
393 $\mathcal{O}((2k + 2^k + 2)2^{3k+1})$.

394 For every such reduced DT skeleton T , we now do the following in order to decide whether
395 T gives rise to a valid record for b . Let $\alpha^{LR \rightarrow} : F_{[k_L]} \cup F_{[k_R]} \rightarrow F_{[k]}$ be the feature relabeling
396 that relabels every (left/right) feature $f_{i_H} \in F_{[k_L]} \cup F_{[k_R]}$ (for some $H \in \{L, R\}$) to its
397 original feature f_i .

398 Let $\alpha^L : F_{[k_R]} \rightarrow I_{[k]}$ be the feature relabeling that relabels every forgotten feature
399 $f_{i_R} \in F_{[k_R]}$ to the future feature $f_{A_{*,i}}$. Let T_L be the reduced DT skeleton obtained from T
400 after applying the relabelling using α^L followed by $\alpha^{LR \rightarrow}$ and then reducing the resulting
401 DT skeleton, i.e., $T_L = r(\alpha^{LR \rightarrow}(\alpha^L(T)))$.

402 Similarly, let $\alpha^R : F_{[k_L]} \rightarrow I_{[k]}$ be the feature relabeling that relabels every forgotten
403 feature $f_{i_L} \in F_{[k_L]}$ to the future feature $f_{A_{i,*}}$. Let T_R be the reduced DT skeleton obtained

404 from T after applying the relabelling using α^R followed by $\alpha^{LR \rightarrow}$ and then reducing the
 405 resulting DT skeleton, i.e., $T_R = r(\alpha^{LR \rightarrow}(\alpha^R(T)))$.

406 Let $\hat{T} = \alpha^{LR \rightarrow}(T)$ and $\hat{s} = |V(T) \setminus V(\hat{T})|$. We now check whether there are records
 407 $(T_L, s_L) \in \mathcal{R}(b_L)$ and $(T_R, s_R) \in \mathcal{R}(b_R)$. If not we discard T and if yes, then we add the
 408 record $(\hat{T}, s_L + s_R + \hat{s})$ to $\mathcal{R}(b)$. This completes the description about how the records
 409 $\mathcal{R}(b)$ are computed. Moreover, the run-time for computing $\mathcal{R}(b)$ can be obtained as follows.
 410 First, because of Lemma 17, we can enumerate all reduced DT skeletons T in time $\mathcal{O}((2k +$
 411 $2^k + 2)2^{3k+1})$. Moreover, computing \hat{T} and \hat{s} can be done in time $\mathcal{O}(|T|) = \mathcal{O}(s)$. Finally,
 412 computing T_L and T_R and checking the existence of the records $(T_L, s_L) \in \mathcal{R}(b_L)$ and
 413 $(T_R, s_R) \in \mathcal{R}(b_R)$ can be achieved in time $\mathcal{O}(?)$. Therefore, we obtain $\mathcal{O}(?)$ as the total
 414 run-time for computing $\mathcal{R}(b)$.

old run-time
 argument below
 should be replaced
 above

415 We now show the correctness of our construction for $\mathcal{R}(b)$, i.e., we have to show that a
 416 record (T, s) is valid if and only if we have added such a record according to our construction
 417 above.

418 Towards showing the forward direction, suppose that (\hat{T}, s) is a valid record in $\mathcal{R}(b)$.
 419 Therefore, there is a DT template T' for b such that $\hat{T} = r(\eta_{\alpha^s_b}(T'))$ and $s = |V(T') \setminus V(T)|$.

420 Because \hat{T} is obtained from T' by reduction, every node in \hat{T} corresponds to a unique
 421 node in T' . Therefore, there is an injective function $z_H : V(\hat{T}) \rightarrow V(T')$ mapping every
 422 node in \hat{T} to its original node in T' . Let T be the DT obtained from \hat{T} after by setting
 423 $\text{feat}_T(t) = i_H$ if $\text{feat}_{\hat{T}}(t) = i$ and $\text{feat}_{T'}(t) \in \text{feat}(b_H)$ for $H \in \{L, B\}$.

424 Note that $\hat{T} = \eta_{\alpha^{LR \rightarrow}}(T)$ and \hat{T} is reduced because $(\hat{T}, s) \in \mathcal{R}(b)$.

425 Let $\alpha^{\rightarrow R} : F_{[k]} \rightarrow F_{[k_R]}$ ($\alpha^{\rightarrow L} : F_{[k]} \rightarrow F_{[k_L]}$) be the feature relabeling that relabels
 426 every forgotten feature $f_i \in F_{[k]}$ to its corresponding forgotten feature in $[k_R]$ ($[k_L]$), i.e.,
 427 $\alpha^{\rightarrow R}(i) = i_R$ ($\alpha^{\rightarrow L}(i) = i_L$) for every $i \in [k]$.

428 Note that $T = r(\eta_{\alpha^{\rightarrow L}}(\eta_{\alpha^s_{b_L}}(\eta_{\alpha^{\rightarrow R}}(\eta_{\alpha^s_{b_R}}(T')))))$.

429 Let $T_L = r(\eta_{\alpha^L}(T))$ and $T_R = r(\eta_{\alpha^R}(T))$. It remains to show that there are s_L and s_R
 430 with $s = s_L + s_R$ such that $(T_L, s_L) \in \mathcal{R}(b_L)$ and $(T_R, s_R) \in \mathcal{R}(b_R)$.

431 Let $T'_L = r(\eta_{\alpha^L}(\eta_{\alpha^{\rightarrow R}}(\eta_{\alpha^s_{b_L}}(T'))))$ and $T'_R = r(\eta_{\alpha^R}(\eta_{\alpha^{\rightarrow L}}(\eta_{\alpha^s_{b_L}}(T'))))$.

432 Note that $T_L = r(\eta_{\alpha^s_{b_L}}(T'_L))$ because of Lemma ?? and the observation that $\eta_{\alpha^s_{b_L}} \circ \eta_{\alpha^L} \circ$
 433 $\eta_{\alpha^{\rightarrow R}} \circ \eta_{\alpha^s_{b_R}} = ?$.

434 Towards showing the reverse direction, suppose that our construction adds the record
 435 $(\hat{T}, s_L + s_R)$ and let T, T_L, T_R be as defined in the construction. Recall that:

- 436 ■ \hat{T} is reduced and $\hat{T} = \eta_{\alpha^{LR \rightarrow}}(T)$,
- 437 ■ $T_L = r(\eta_{\alpha^L}(T))$ and $(T_L, s_L) \in \mathcal{R}(b_L)$,
- 438 ■ $T_R = r(\eta_{\alpha^R}(T))$ and $(T_R, s_R) \in \mathcal{R}(b_R)$.

439 Let T'_L be the reduced DT template for b_L such that $T_L = r(\eta_{\alpha^s_{b_L}}(T'_L))$ and $s_L =$
 440 $|V(T'_L) \setminus V(T_L)|$, which exists because $(T_L, s_L) \in \mathcal{R}(b_L)$. Similarly, let T'_R be the reduced
 441 DT template for b_R such that $T_R = r(\eta_{\alpha^s_{b_R}}(T'_R))$ and $s_R = |V(T'_R) \setminus V(T_R)|$, which exists
 442 because $(T_R, s_R) \in \mathcal{R}(b_R)$.

443 We now show how to construct a witness T' (from T, T'_L , and T'_R) for the validity of the
 444 record $(\hat{T}, s_L + s_R)$, i.e., T' is a reduced DT template for b such that $\hat{T} = r(\alpha^s_b(T'))$ and
 445 $s_L + s_R = |V(T') \setminus V(\hat{T})|$.

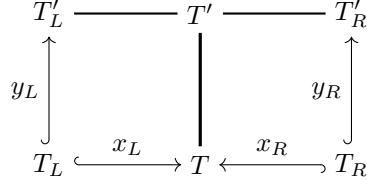
todo:show
 minimality here
 maybe it can be
 done using the
 forward direction

446 Suppose that there is a reduced DT template T' for b such that $\hat{T} = r(\alpha^s_b(T'))$ and
 447 $|V(T') \setminus V(\hat{T})| < s_L + s_R$.

448 Informally, we obtain T' from T after reversing the relabelling and reduction operations
 449 applied to T'_L and T'_R to obtain T_L and T_R , respectively; recall that $T_H = r(\eta_{\alpha^s_{b_H}}(T'_H))$ for

450 $H \in \{L, R\}$. That is, we will reverse the labelling for the nodes in T and add back the nodes
 451 to T that have been removed from T'_L and T'_R .

452 Let $H \in \{L, R\}$. Because T_H is obtained from T by reduction, every node in T_H
 453 corresponds to a unique node in T . Therefore, there is an injective function $x_H : V(T_H) \rightarrow$
 454 $V(T)$ mapping every node in T_H to its original node in T . Similarly, because T_H is obtained
 455 from T'_H by reduction, there is an injective function $y_H : V(T_H) \rightarrow V(T'_H)$ mapping every
 456 node in T_H to its original node in T'_H . See also Figure 2 for an illustration of these mappings.



■ **Figure 2**

457 Our first order of business is to rename all forgotten features in T to their real features
 458 as given by T'_L and T'_R . That is, for every node t in T assigned to a forgotten feature, i.e.,
 459 $\text{feat}(t) \in F_{[k_L]} \cup F_{[k_R]}$, we do the following. If $\text{feat}(t) \in F_{[k_H]}$ for $H \in \{L, R\}$, then t is also
 460 in T_H and hence also in T'_H . Therefore, we can change $\text{feat}(t)$ to the real feature assigned to
 461 t in T'_H . Let T^0 be the DT obtained from T after renaming all forgotten features to real
 462 features in this manner.

todo: explain

463 Consider an edge $e = (p, c)$ in T_L such that p is the parent of c in T_L . Then, e corresponds
 464 to a path $P'_L(e)$ between $y_L(p)$ and $y_L(c)$ in T'_L . Similarly, e corresponds to a path $P_L(e)$
 465 between $x_L(p)$ and $x_L(c)$ in T^0 .

466 Our next order of business is now to add all nodes to T^0 that have been removed when
 467 going from T'_L to T_L (via the reduction $r(\eta_{\alpha_{T'_L}^s}(T'_L))$). To achieve this, we go over every edge
 468 $e = (p, c)$ of T_L such that p is the parent of c in T_L and plugin the path $P'_L(e)$ (from T'_L)
 469 into the last edge on the path $P_L(e)$ (from T^0). Let T^1 be the tree obtained from T^0 after
 470 doing this operation for every edge of T_L .

471 Consider an edge $e = (p, c)$ in T_R such that p is the parent of c in T_R . Then, e corresponds
 472 to a path $P'_R(e)$ between $y_R(p)$ and $y_R(c)$ in T'_R . Similarly, e corresponds to a path $P_R(e)$
 473 between $x_R(p)$ and $x_R(c)$ in T^1 . Similarly to above, we now add all nodes to T^1 that have
 474 been removed when going from T'_R to T_R (via the reduction $r(\eta_{\alpha_{T'_R}^s}(T'_R))$). To achieve this,
 475 we go over every edge $e = (p, c)$ of T_R such that p is the parent of c in T_R and plugin the
 476 path $P'_R(e)$ (from T'_R) into the last edge on the path $P_R(e)$ (from T^1). Let T' be the tree
 477 obtained from T^1 after doing this operation for every edge of T_R .

478 We now show that T' is indeed a witness for the validity of the record $(\hat{T}, s_L + s_R)$, i.e.,
 479 T' is a reduced DT template for b such that $\hat{T} = r(\alpha_b^s(T'))$ and $s_L + s_R = |V(T') \setminus V(\hat{T})|$.

480 We start by showing that $\hat{T} = r(\eta_{\alpha_b^s}(T'))$. Because $\hat{T} = \alpha_b^s(T^0)$, it suffices to show that
 481 the only nodes removed from T' are the ones that we added to T^0 to obtain T' . Or in other
 482 words, we need to show that only the nodes that are redundant in $\eta_{\alpha_b^s}(T')$ are the nodes in
 483 $V(T') \setminus V(T^0)$.

484 Consider a node $t \in V(T') \setminus V(T^0)$, i.e., t is a node that we added to T^0 to obtain T' .
 485 Then, $t \in V(T'_H) \setminus V(T_H)$ for some $H \in \{L, R\}$. Because $T_H = r(\eta_{\alpha_{b_H}^s}(T'_H))$, t is redundant
 486 in $\eta_{\alpha_{b_H}^s}(T'_H)$, because of some node $t' \in V(T_H)$ with $\alpha_{b_H}^s(\text{feat}_{T'_H}(t)) = \alpha_{b_H}^s(\text{feat}_{T'_H}(t'))$. Since
 487 $t' \in V(T_H)$ also $t' \in V(T')$ and therefore t is also redundant in $\eta_{\alpha_b^s}(T')$ (because of t'), as
 488 required.

489 Now consider a node $t \in V(T^0)$ and assume for a contradiction that t is redundant in
 490 $\alpha_b^s(T')$ because of some node $t' \in V(T')$ with $\alpha_b^s(\text{feat}_{T'}(t)) = \alpha_b^s(\text{feat}_{T'}(t'))$. Then, because
 491 $\hat{T} = \alpha_b^s(T^0)$ is reduced, we obtain that $t' \in V(T') \setminus V(T^0)$. Therefore, $t' \in V(T'_H) \setminus V(T_H)$
 492 for some $H \in \{L, R\}$. But then, t' is redundant in $\eta_{\alpha_b^s}(T'_H)$ because of some node $t'' \in V(T_H)$
 493 with $\alpha_b^s(\text{feat}_{T'}(t'')) = \alpha_b^s(\text{feat}_{T'_H}(t'))$, which implies that also t is redundant in \hat{T} because of
 494 t'' a contradiction to our assumption that \hat{T} is reduced. This shows that $\hat{T} = r(\eta_{\alpha_b^s}(T'))$.
 495 Moreover, because $|V(T^0)| = |V(\hat{T})|$ and $|V(T') \setminus V(T^0)| = s_L + s_R$, it also follows that
 496 $s_L + s_R = |V(T') \setminus V(\hat{T})|$.

proof of
minimality of T' 497
still missing 498

496 Moreover, $V(T) \setminus Im(x_H)$ and $V(T'_H) \setminus Im(y_H)$ can be partitioned into subtrees that
 have been deleted after the application of $r \circ p_*$, $r \circ p'_*$ on T or of the standard reduction
 499 on T'_H : let X_H^* and Y_H^* be the set of roots of the above subtrees in $V(T) \setminus Im(x_H)$ and
 500 $V(T'_H) \setminus Im(y_H)$ respectively. In addition, for every element $y \in Y_H^*$, let Y_y^H be the maximal
 501 subtree of T'_H rooted at y with no elements from $Im(y_H)$ and that does not contain any
 502 vertex from $Y_H^* \setminus \{y\}$; let (Y_y^H, S_y^H) the corresponding single pair. In a similar way, for every
 503 element $x \in X_H^*$, let X_x^H be the maximal subtree of T rooted at x with no elements from
 504 $Im(x_H)$ and that does not contain any vertex from $X_H^* \setminus \{x\}$; let (X_x^H, S_x^H) the corresponding
 505 single pair. Finally, for every $y \in Y_H^*$, let P_y^H be the shortest downwards path in T'_H that
 506 contains y and with both endpoints in $Im(y_H)$, say $y_H(t)$ and $y_H(t')$.

507 *Claim 1:* For every $H \in \{L, R\}$ and for every $y, y' \in Y_H^*$, the paths P_y^H and $P_{y'}^H$ are either
 508 edge disjoint or $P_y^H = P_{y'}^H$.

509 *Proof.* If P_y^H and $P_{y'}^H$ are edge disjoint, then the statement is proven immediately. Suppose
 510 P_y^H and $P_{y'}^H$ share an edge. By minimality and the fact they are downwards paths, P_y^H and
 511 $P_{y'}^H$ share the endpoint towards the root. If they also share the other endpoint, then the
 512 statement is proven immediately. Suppose now their endpoints towards the leaves is different,
 513 say w and w' , and consider the last edge those paths have in common in a root-to-leaf order,
 514 say uv .

515 Without loss of generality, we can assume w belongs to the left branch of v and w' belongs
 516 to the right branch of v . Note that $v \in V(T'_H) \setminus Im(y_H)$, or we get a contradiction due the
 517 minimality of P_y^H . Now we get the following contradiction: by construction, w and w' are
 518 both elements of $Im(y_H)$ but at least one of them must be in $V(T'_H) \setminus Im(y_H)$ since it is an
 519 element of either Y_y^H or of $Y_{y'}^H$. This proves Claim 1.

520 Now for every $y \in Y_H^*$ we consider the path Q_y^H in T having endpoints $x_H(t)$ and $x_H(t')$.

521 Now we are able to describe how to obtain a witness T' of T for b . For every $y \in Y_L^*$, in
 522 the last edge of path Q_y^L we plug in the single pair $(Y_{y'}^L, S_{y'}^L)$ rooted at y' , for every internal
 523 node y' of P_y^L , in the order the nodes y' appear in P_y^L . Note that, in the case an element
 524 of Y_L^* is present in more than one P_y^L , we plug in the corresponding single pair only once.
 525 Note also that whenever we plug in some single pair (Y_y^L, S_y^L) in a DT, the tree Y_y^L has real
 526 features and future features as nodes. Call this graph T^* . Now we do the same sequence of
 527 plug ins of the single pairs corresponding to the internal vertices of P_y^R in the last edge of
 528 the path Q_y^R . Again, in the case an element of Y_R^* is present in more than one P_y^R , we plug
 529 in the corresponding single pair only once. Call the tree obtained in this way T' . Node that
 530 T' contains real features from $\text{feat}(b_L)$ and from $\text{feat}(b_R)$ and future features with labels in
 531 $\mathcal{P}([k])$.

532 To conclude this part of the proof we have to show two things: (i) T is obtained from T'
 533 after removing s vertices; (ii) T' is a real DT for b . We start proving (i): by construction T'
 534 is obtained from T after adding s_L elements from T'_L and s_R elements from T'_R , and so with
 535 $s_L + s_R = s$ more elements.

536 Before considering statement (ii), we consider the following relabelling p_+ of T' : every
 537 real feature in $feat(b_R)$ is assigned to a feature with its label at node b_R and every other
 538 feature is assigned to itself. The real DT T'_L can be obtained from T' by the application of
 539 the composition $r \circ p_* \circ p_+$.

540 Now we consider statement (ii). We show that given an example $e \in exam(b_L)$, e is
 541 correctly classified by T' and to do so we show that e ends in a leaf of T' that corresponds
 542 to the leaf where e ends in T'_L . Say that e goes along a path P of T'_L from the root to a
 543 leaf ℓ and let Q be the corresponding path in T' , i.e. the path from r to ℓ (note that by
 544 construction ℓ is present in T' and is still a leaf). Let v be a node of Q , we can have the
 545 following different cases.

- 546 ■ v is a real feature from $feat(b_L)$: v is also present in T'_L as real feature;
- 547 ■ v is a real feature from $feat(b_R)$: v might not be present in T'_L due reductions but if it is
 present it is a future feature A_i for some $i \in [k]$;
- 549 ■ v is a future feature f_A : v might not be present in T'_L due reductions but if it is present
 it is still the same future feature A_i .

551 If v is present in T'_L then the behaviour of v on e in T'_L and in T' is the same. Suppose
 552 now v is a node of Q that is being reduced due his label and so it is not present in T'_L .
 553 This means there is a set of ancestors of v such that their labels allows to remove v and by
 554 construction v behaves on e like those ancestors. This proves e goes along Q and in particular
 555 it ends at leaf ℓ and so T' is a real DT for b_L . With symmetric construction, we show that
 556 T' is also a real DT for b_R .

557 Now we prove the backward direction. Let T be a reduced DT such that s is the minimum
 558 number of elements that have been deleted from a witness T' of T for b . In particular, we
 559 recall that T' is a real DT for b with actual feature labels in $[k] \cup [k']$ and future feature
 560 labels in $\mathcal{P}([k])$.

561 We create at real DT T'_L by the application of the composition $r \circ p_* \circ p_+$ to T' . By
 562 assumption T' is a real DT for b_L and by construction T'_L is a real DT for b_L . Denote
 563 with T_L the DT template obtained from T'_L by standard reduction and denote with s_L
 564 the number of nodes that have been deleted from T'_L to obtain T . By induction we have
 565 $(T_L, s_L) \in \mathcal{R}(b_L)$. Now we note that T_L is obtained from T after the application of the
 566 composition $r \circ p_*$. In a symmetric way, we construct T'_R , T_R and the record $(T_R, s_R) \in \mathcal{R}(b_R)$.
 567 Then $(T, s_L + s_R) \in \mathcal{R}(b)$. ◀

568 « « « < HEAD

569 ▶ **Lemma 14** (relabel node). *Let $b \in V(B)$ be relabel node. Then $\mathcal{R}(b)$ can be computed in
 570 time $\mathcal{O}(k(2k + 2^k + 2)2^{3k+1})$.*

571 **Proof.** Let b_C be the unique child of b in B . Let R be the mapping of $[k]$ to itself that
 572 represent the node b . Moreover, since we are considering a nice NLC-expression we can
 573 assume R is the identity mapping, i.e. $R(\ell) = \ell$, for all values except for a unique element i
 574 of its domain, i.e. $R(i) = j$ for some $j \in [k] \setminus \{i\}$.

575 We say that a future feature A is *good* if it does not distinguish between i and j , that
 576 is $i \in A$ if and only if $j \in A$, and *bad* otherwise. Let (T_C, s_C) be an element of $\mathcal{R}(b_C)$. Let
 577 p'' the following relabelling of the DT template T_C : every feature with label i is assigned
 578 to label j and every future feature with label A is assigned to the future feature with label
 579 $A \setminus \{i\}$.

580 If T_C has a bad future feature then we do not take any other action. Suppose now T_C
 581 has only good future features; now let T be the DT template obtained from T_C after the
 582 application of the composition $r \circ p''$ and let s^* be the number of nodes that have been
 583 deleted from T_C to T .

584 If there is a record in $\mathcal{R}(b)$ of the form (T, s') for some integer $s' \leq s_C + s^*$ then we do
 585 not take any other action. If there is a record in $\mathcal{R}(b)$ of the form (T, s') for some integer
 586 $s' > s_C + s^*$ then we replace it with $(T, s_C + s^*)$. If there is no record in $\mathcal{R}(b)$ of the form
 587 (T, s') for some integer s' then we add $(T, s_C + s^*)$ to $\mathcal{R}(b)$.

588 Now we want to evaluate the running time of computing $\mathcal{R}(b)$. Consider record (T_C, s_C)
 589 in $\mathcal{R}(b_C)$. In $\mathcal{O}(k)$ time we check if T_C all the future features are good. For every such DT
 590 T_C , there are at most 2^{2k} paths from the root to the leaves and for every of these paths there
 591 are at most k nodes for each of the following: feature with label i and and future feature
 592 that contains i . This means $r \circ p''$ can be done in $\mathcal{O}(k)$ time. This means to compute $\mathcal{R}(b)$
 593 takes $\mathcal{O}(k|\mathcal{R}(b_C)|) = \mathcal{O}(k(2k + 2^k + 2)2^{3k+1})$ time.

594 Now we have to show the correctness of the construction for $\mathcal{R}(b)$, i.e. $(T, s) \in \mathcal{R}(b)$ if
 595 and only if s is the minimum number of elements that have been deleted from a witness T'
 596 of T for b .

597 We start with the forward direction. Let $(T, s) \in \mathcal{R}(b)$. By construction there exists a
 598 record $(T_C, s_C) \in \mathcal{R}(b_C)$ such that T is obtained from T_C after the application of $r \circ p''$ and
 599 let $s^* = s - s_C$. By induction s_C is the minimum amount of nodes that have been deleted
 600 from a witness T'_C of T_C for b_C . By construction we also know that every future feature of
 601 both T'_C and T_C is good.

602 Denote with T' the real DT obtained T'_C after the application of $r \circ p''$: note that this
 603 last reduction does not any node since every future feature of T'_C is good and there is no
 604 feature with label i . To conclude this part of the proof we have to show two things: (i) T is
 605 obtained from T' after removing s vertices; (ii) T' is a witness of T for b .

606 Before proving (i), we describe how T can be obtained from T' . Let p''' be the following
 607 relabelling of T' : every real feature that contains j is assigned to the real feature $A \cup \{i\}$
 608 and every other feature is assigned to itself. Then the application of the composition p''' ,
 609 the standard reduction and $r \circ p''$ to T' is exactly the standard reduction for T' which then
 610 result to the DT template T . By Lemma 15 the score of the standard reduction from T' to
 611 T is exactly $s_C + s^* = s$.

612 Now we consider statement (ii). First note that $\text{exam}(b) = \text{exam}(b_C)$. We show that
 613 a given example $e \in \text{exam}(b)$ is correctly classified by T' . Say that e goes along a path P
 614 of T'_C from the root to a leaf ℓ . We show e goes along the path P in T' as well: every real
 615 feature has not changed and so e behaves the same. Since every future feature of T'_C is good,
 616 then e behave the same on the corresponding future feature of T' .

617 Now we prove the backward direction. Let T be a reduced DT such that s is the minimum
 618 number of elements that have been deleted from a witness T' of B for b . In particular, we
 619 recall that real T' is a DT for b with real features and future feature labels in $\mathcal{P}([k] \setminus \{i\})$.

620 We create the real DT T'_C as the application of $r \circ p'''$ to T' , the DT template T_C as the
 621 application of the standard reduction to T'_C . By construction we have $(T_C, s_C) \in \mathcal{R}(b_C)$,
 622 where s_C is the number of nodes that have been removed from T'_C to T_C . Note that T_C has
 623 only good future features. Finally we note that T is obtained from T_C by the application of
 624 $r \circ p''$. ◀

625 3.4 Formal Definition of Records and Preliminary Results

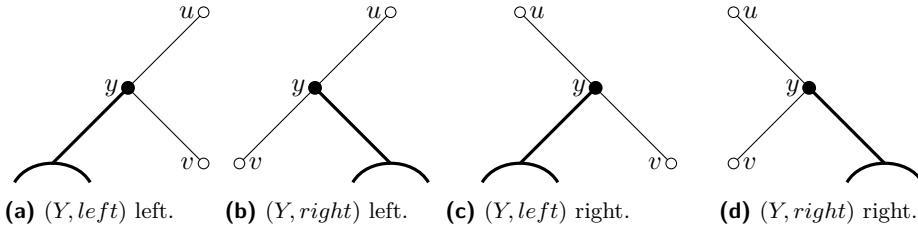
626 =====



627 **NLC-width**

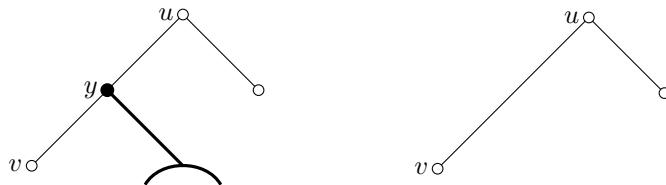
628 »»»> e150fdde332112fd1c2acb6bd85a9a5606b79547 We start off with some definitions. We
 629 say an edge is a *left* (*right*) *edge* of a subcubic rooted tree if it connects a non-leaf node with
 630 his left (resp. right) child. Let Y be a rooted subcubic tree and $S \in \{\text{left}, \text{right}\}$, then we
 631 say the pair (Y, S) is a *single pair* if the root of Y has at most one child and the side S
 632 indicates whether the edge from the root is either a left or right edge. Moreover, we say that
 633 (Y, S) is single pair in a subcubic rooted tree T if Y is a maximal subtree of T and in Y the
 634 root have at most the S child. Note that when tree of a single pair is made of just a node,
 635 the side is not relevant.

636 Now we can define two operations on subcubic rooted trees and single pairs. We say that
 637 we *plug in* a single pair (Y, S) in a left (right) edge uv as follows: we make the root y of Y the
 638 left (right) child of u , $Y \setminus \{y\}$ to be the S subtree of y and v to be the $H \in \{\text{left}, \text{right}\} \setminus S$
 639 child of y . See Figure 3 for the corresponding drawings. Note after a plug in of a single pair
 640 in an edge, the node v belongs in the same side of the subtree rooted at u as it was before
 641 the plug in.



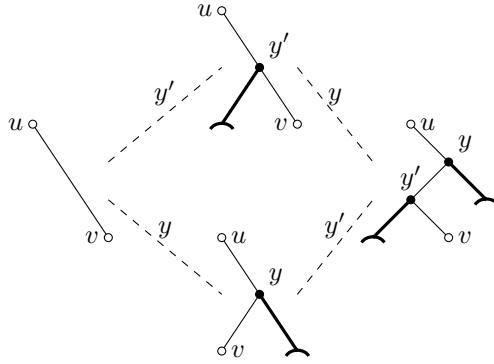
■ **Figure 3** The drawings describe the plug in operation in the different four cases. The bold part highlight the single pair (Y, S) .

642 Let (Y, S) be a single pair in a rooted subcubic tree T , then we *remove* (Y, S) from T as
 643 follows. Let y be the root of Y . If y is the root of T , then we obtain an empty tree. If y is a
 644 leaf node of T , then we obtain $T - y$. Otherwise let y be a non-root and non-leaf node, let u
 645 be the parent of y and v be the child of y that is not in $V(Y)$, then we consider the tree
 646 obtained from T after replacing y with v as the child of u and deleting Y . See Figure 4 for
 647 an example.



■ **Figure 4** The drawing describe an example of the remove operation: a single pair (Y, right) is removed from a subcubic rooted tree. The bold part highlight the single pair (Y, S) .

648 It is clear from the four different plug in cases that if we want to plug in two pairs (Y, S)
 649 and (Y', S') on an edge uv such that the ancestor-descendant relationship is given, say y of
 650 Y has to be in the path from the root to y' of Y' , then we can do these plug ins in any order
 651 but with some care. It is the same if we first plug in (Y, S) in the edge uv and then plug in
 652 (Y', S') in the edge vy or if we first plug in (Y', S') in the edge uv and then plug in (Y, S) in
 653 the edge uy' . See Figure 5 for the an example.



■ **Figure 5** An example of plugging in two pairs (Y, left) and (Y', right) in a left edge uv .

654 For a subset of labels $A \subseteq [k]$, we define the feature template f_A by setting $e(f_A) = 1$ if
 655 and only if $\text{lab}(e) \in A$ and $e(f_A) = 0$ otherwise. With a small abuse of notation, we often
 656 identify the feature template f_A with the corresponding subset of labels A .

657 Suppose we have a DT such that some feature label i occurs twice on a path from the
 658 root to the leaves, say f_1 is the instance closer to the root and f_2 is the other instance. If f_2
 659 is in the left (resp. right) subtree of f_1 , we remove f_2 's right (resp. left) subtree. In this case
 660 we say we have done an *actual removal*.

661 Suppose we have a feature template labelled A in our decision tree. Let A_1, \dots, A_ℓ be the
 662 sequence of feature templates on the path from the root to A in order (not including A). Let
 663 $A'_i = A_i$ if A is in the right sub-tree of A_i and let $A'_i = \overline{A_i}$ otherwise. If $\overline{A} \subseteq A'_1 \cup \dots \cup A'_\ell$,
 664 then we remove the subtree rooted at the left child of A . If $A \subseteq \overline{A'_1} \cup \dots \cup \overline{A'_\ell}$, then we
 665 remove the subtree rooted at the right child of A . In this case we say we have done a *template*
 666 *removal*. If this procedure has been applied to a record exhaustively, we say that the DT is
 667 *reduced*.

668 To be short, for a DT T and a node v , we write $v \in T$ instead of $v \in V(T)$ and $v \notin T$
 669 otherwise. In a DT T we say that path p is a *downward* path if it is contained in a
 670 path having the root as endpoint.

671 We now formally define two important operations. Given a DT T , we say that we *reduce*
 672 T if we exhaustively do actual removals and template removals. Call $r(T)$ the resulting DT.

673 Recall that in any DT T , every non-leaf node v has one of the following three contents: v
 674 is a real feature (without label), or v is a feature with a label, or v is a future feature with
 675 the corresponding subset of labels. A *relabelling* p for T is an assignment of contents of T
 676 as follows. Every feature is assigned to a feature with either future, real or with a label.
 677 We say that we *relabel* the DT T via the relabelling p if for every node of T we apply the
 678 corresponding assignment and call $p(T)$ the resulting DT.

679 The following lemma shows that, after repeatedly applying it the necessary amount of
 680 times, to obtain a reduced DT after a sequence of relabels, it is safe to reduce at the end.

681 ▶ **Lemma 15** (Relabelling Lemma). *Let T be a DT and p be relabelling of T . Then $(r \circ p \circ r)(T) = (r \circ p)(T)$.*

683 **Proof.** For every $v \in T$, we want to prove $v \in (r \circ p \circ r)(T) \Leftrightarrow v \in (r \circ p)(T)$.

684 \Rightarrow Suppose there is a node $v \notin (r \circ p)(T)$. Since $v \in p(T)$, there is a set of ancestors of v
 685 in $p(T)$ that allows to remove v . Let A_v be the union of all the minimal set of ancestors of v
 686 in $p(T)$ that allows to remove v . If A_v is a set of ancestors of v in T that allows to reduce v

687 then $v \notin r(T)$ and so $v \notin (r \circ p \circ r)(T)$. Otherwise let A'_v be the subset of A_v in $(p \circ r)(T)$.
 688 We conclude by noting that A'_v contains one of the minimal sets A_v is composed of and so
 689 $v \notin (r \circ p \circ r)(T)$.

690 \Leftarrow Suppose there is a node $v \notin (r \circ p \circ r)(T)$. If $v \in (p \circ r)(T)$, there exists a set A_v of
 691 ancestors of v in $(p \circ r)(T)$ that allows to reduce v . Then A_v is a set of ancestors of v in $p(T)$
 692 that allows to reduce v and so $v \notin (r \circ p)(T)$. If $v \notin (p \circ r)(T)$ then $v \notin r(T)$: there exists a
 693 set A_v of ancestors of v in T that allows to remove v . This means A_v is a set of ancestors of
 694 v in $p(T)$ that allows to remove v and so $v \notin (r \circ p)(T)$. \blacktriangleleft

695 We say that a DT T is a *real DT* if every non-leaf node is either a real feature or a future
 696 feature, whereas it is a *DT template* if it contains no real feature.

697 Let B be a rooted subcubic tree that corresponds to a k -NLC expression of the graph
 698 $G_I(E)$. For $b \in V(B)$, we write $feat(b)$ and $exam(b)$ for the sets of features and examples
 699 introduced at node b . We say that a real DT T is a DT for the node b if every real feature of
 700 T is an element of $feat(b)$ and every example in $exam(b)$ is correctly classified by T , i.e. if
 701 $e \in exam(b) \cap E^+$ then e ends in a leaf with a + label and if $e \in exam(b) \cap E^-$ then e ends
 702 in a leaf with a - label.

703 Given a real DT T and a node $b \in B$, often we want to perform a very specific composition
 704 of operations. Let p_b be the following relabelling of T : every real feature of T is assigned to
 705 a feature with the label given by the k -NLC expression at node b and every other feature is
 706 assigned to itself. Then the composition $r \circ p_b$ is called the *standard reduction* of T at node
 707 b . Given a DT T and a node $b \in B$, it is useful to give the following relabelling p'_b : every
 708 feature with a label is assigned to the real feature of that node. The relabelling p'_b is called
 709 the *real relabelling* of T at node b .

710 We say that a DT template T is a DT for the node b if there exists a real DT T' for b such
 711 that T is the standard reduction of T' . In this case we say that T' is the witness of T for b .

712 \blacktriangleright **Lemma 16.** *If there are ℓ features with labels and 2^h future features, then every reduced
 713 DT template has height at most $\ell + h$. Furthermore, every path from the root to the leaves
 714 contains at most ℓ features with label and at most $h - 1$ future features.*

715 **Proof.** Consider a path P of maximum length from the root to the leaves in a reduced DT
 716 template T . By the assumptions on T , no feature with label appears more than once on
 717 this path: the number of these feature nodes on this path is at most ℓ . Consider two future
 718 features f_A and $f_{A'}$ that appear in P , say f_A is the instance closer to the root. Since T is
 719 reduced, we must have that $\emptyset \subset A' \subset A$. Since the label of any future feature has at most h
 720 elements, there can be at most $h - 1$ feature template nodes on this path. The path ends
 721 with a leaf node, so this gives a total of $\ell + h - 1 + 1 = \ell + h$ nodes, as required. \blacktriangleleft

722 \blacktriangleright **Lemma 17.** *If there are ℓ features with label and 2^h future features, then there are at
 723 most $(\ell + 2^k + 2)2^{\ell+k+1}$ reduced DT templates. Furthermore, these can be enumerated in
 724 $\mathcal{O}((\ell + 2^k + 2)2^{\ell+k+1})$ -time.*

725 **Proof.** By Lemma 16, the tree has height at most $\ell + k$. Each node of the decision tree could
 726 be a feature with label, a future feature, or a leaf: at most $\ell + 2^h + 2$ different contents. Since
 727 there are at most $2^{\ell+h+1}$ nodes in the tree, there are at most $(\ell + 2^h + 2)2^{\ell+h+1}$ possible
 728 decision trees. \blacktriangleleft

729 The *semantics* for a record are defined as follows. We say that a pair (T, s) is a *record* for
 730 the node $b \in B$ and we write $(T, s) \in \mathcal{R}(b)$, if T is a DT template for b and s is the minimum
 731 number of elements that have been deleted from a witness T' of T for b .

732 3.5 Proof to the Main Result

733 Now, it suffices to compute $\mathcal{R}(b)$ via leaf-to-root dynamic programming. The following
 734 four lemmas show how this can be achieved for all of the four types of nodes in a k -NLC
 735 expression tree B .

736 ▶ **Lemma 18** (leaf node). *Let $b \in V(B)$ be a leaf node. Then $\mathcal{R}(b)$ can be computed in time
 737 $\mathcal{O}(k(2^k + 3)2^{k+2})$.*

738 **Proof.** Let v be the vertex of $G_I(E)$ that corresponds to the leaf node b . This means either
 739 $v \in E$ or $v \in \text{feat}(E)$.

740 We have to enumerate all possible reduced DT templates T for b . It is enough to consider
 741 all reduced DT templates T of height at most $k + 1$ and discard those that are not DT
 742 templates for b ; these can be enumerated in time $\mathcal{O}((2^k + 3)2^{k+2})$ by Lemma 17 and the
 743 check can be done in time $\mathcal{O}(k)$. We add the pair $(T, 0)$ to the set of records $\mathcal{R}(b)$.

744 Now we have to show the correctness of the construction for $\mathcal{R}(b)$, i.e. $(T, s) \in \mathcal{R}(b)$ if
 745 and only if s is the minimum number of elements that have been deleted from a witness T'
 746 of T for b .

747 We start with the forward direction. Let $(T, s) \in \mathcal{R}(b)$. By construction, we have that
 748 $s = 0$ and T is a DT template for b which is already reduced. Then T is trivially a witness
 749 of T for b .

750 Now we prove the backward direction. Let T be a reduced DT template such that 0
 751 is the minimum number of elements that have been deleted from a witness T' of T for b .
 752 This means T' is obtained from T after the real relabelling at node b is applied: T is a DT
 753 template among the considered DTs above which leads to the fact that $(T, 0) \in \mathcal{R}(b)$. ◀

754 ▶ **Lemma 19** (join node). *Let $b \in V(B)$ be a join node. Then $\mathcal{R}(b)$ can be computed in time
 755 $\mathcal{O}(k(2k + 2^k + 2)2^{6k+1})$.*

756 **Proof.** Let b_L and b_R be the left, resp. right, child of b in B : we may assume the labels for
 757 $\text{feat}(b_L)$ are in $[k]$ and the labels for $\text{feat}(b_R)$ are in $[k']$. Moreover, let M be the $k \times k$ $\{0, 1\}$
 758 matrix that represent the node b . Finally, for every label $i \in [k]$, let $A_i = \{j \in [k] \mid M_{i,j} = 1\}$.

759 We consider every reduced DT T for b with feature labels in $[k] \cup [k']$ and future feature
 760 labels in $\mathcal{P}([k])$; these can be enumerated in time $\mathcal{O}((2k + 2^k + 2)2^{3k+1})$ by Lemma 17.

761 For every such DT T , we create a DT T_L as follows. Let p_* be the following relabelling:
 762 for every $i' \in [k']$, every feature with label i' is assigned to the future feature A_i . Then we
 763 apply the composition $r \circ p_*$ to T . In a symmetrical way we create a DT T_R . Let p'_* be the
 764 following relabelling: for every $i \in [k]$, every feature with label i is assigned to the future
 765 feature $A_{i'}$ and every future feature A_i is assigned to the future feature $A_{i'}$. Then we apply
 766 the composition $r \circ p'_*$ to T .

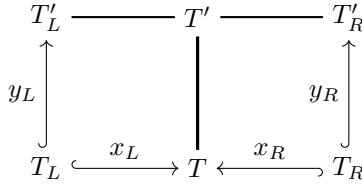
767 Now we want to understand if there is a record in $\mathcal{R}(b_L)$ of the form (T_L, s_L) for some
 768 positive integer s_L and if there is a record in $\mathcal{R}(b_R)$ of the form (T_R, s_R) for some positive
 769 integer s_R : if the answer is yes in both cases, we add a record $(T, s_L + s_R)$ to $\mathcal{R}(b)$; otherwise
 770 we discard this option.

771 Now we want to evaluate the running time of computing $\mathcal{R}(b)$. Every reduced DT T can
 772 be enumerated in time $\mathcal{O}((2k + 2^k + 2)2^{3k+1})$ by Lemma 17. For every such DT T , there are
 773 at most 2^{3k} paths from the root to the leaves and for every of these paths there are at most
 774 k nodes for each of the following: features with label in $[k]$, features with label in $[k']$ and
 775 future features by Lemma 16. This means $r \circ p_*$ and $r \circ p'_*$ can be done in $\mathcal{O}(k2^{3k})$ time.

776 Now we have to show the correctness of the construction for $\mathcal{R}(b)$. We start with the
 777 forward direction. Let $(T, s) \in \mathcal{R}(b)$. By construction there exist records $(T_L, s_L) \in \mathcal{R}(b_L)$
 778 and $(T_R, s_R) \in \mathcal{R}(b_R)$ such that T_L and T_R are obtained by the application of $r \circ p_*$ and
 779 $r \circ p'_*$ respectively to T and $s_L + s_R = s$.

780 By induction, for $H \in \{L, R\}$, we know that s_H is the minimum number of elements that
 781 have been deleted from a witness T'_H of T_H for b_H .

782 For $H \in \{L, R\}$, we define maps x_H and y_H as follows. Let $x_H : V(T_H) \rightarrow V(T)$ and
 783 $y_H : V(T_H) \rightarrow V(T'_L)$ be the functions that maps every node of T_H to the corresponding
 784 node in T and in T'_L and note that by constructions both these maps are injective.



785 Moreover, $V(T) \setminus Im(x_H)$ and $V(T'_H) \setminus Im(y_H)$ can be partitioned into subtrees that
 786 have been deleted after the application of $r \circ p_*$, $r \circ p'_*$ on T or of the standard reduction
 787 on T'_H : let X_H^* and Y_H^* be the set of roots of the above subtrees in $V(T) \setminus Im(x_H)$ and
 788 $V(T'_H) \setminus Im(y_H)$ respectively. In addition, for every element $y \in Y_H^*$, let Y_y^H be the maximal
 789 subtree of T'_H rooted at y with no elements from $Im(y_H)$ and that does not contain any vertex
 790 from $Y_H^* \setminus \{y\}$; let (Y_y^H, S_y^H) the corresponding single pair. <<<< HEAD ===== >>>
 791 e150fdde332112fd1c2acb6bd85a9a5606b79547 In a similar way, for every element $x \in X_H^*$, let
 792 X_x^H be the maximal subtree of T rooted at x with no elements from $Im(x_H)$ and that does
 793 not contain any vertex from $X_H^* \setminus \{x\}$; let (X_x^H, S_x^H) the corresponding single pair. Finally,
 794 for every $y \in Y_H^*$, let P_y^H be the shortest downwards path in T'_H that contains y and with
 795 both endpoints in $Im(y_H)$, say $y_H(t)$ and $y_H(t')$.

796 *Claim 1:* For every $H \in \{L, R\}$ and for every $y, y' \in Y_H^*$, the paths P_y^H and $P_{y'}^H$ are either
 797 edge disjoint or $P_y^H = P_{y'}^H$.

798 *Proof.* If P_y^H and $P_{y'}^H$ are edge disjoint, then the statement is proven immediately. Suppose
 799 P_y^H and $P_{y'}^H$ share an edge. By minimality and the fact they are downwards paths, P_y^H and
 800 $P_{y'}^H$ share the endpoint towards the root. If they also share the other endpoint, then the
 801 statement is proven immediately. Suppose now their endpoints towards the leaves is different,
 802 say w and w' , and consider the last edge those paths have in common in a root-to-leaf order,
 803 say uv .

804 Without loss of generality, we can assume w belongs to the left branch of v and w' belongs
 805 to the right branch of v . Note that $v \in V(T'_H) \setminus Im(y_H)$, or we get a contradiction due the
 806 minimality of P_y^H . Now we get the following contradiction: by construction, w and w' are
 807 both elements of $Im(y_H)$ but at least one of them must be in $V(T'_H) \setminus Im(y_H)$ since it is an
 808 element of either Y_y^H or of $Y_{y'}^H$. This proves Claim 1.

809 Now for every $y \in Y_H^*$ we consider the path Q_y^H in T having endpoints $x_H(t)$ and $x_H(t')$.

810 <<<< HEAD ===== *Claim 2:* For every $H \in \{L, R\}$ and for every $y \in Y_H^*$, every
 811 internal vertex of Q_y^H is an element of X_H^* .

812 *Proof.* Suppose that Q_y^H has an internal vertex $t \notin X_H^*$. By definition, there exists a vertex
 813 $v \in V(T_H)$ such that $x_H(v) = t$. Since x_H is injective then $v \notin \{v_1, v_2\}$. Since y_H is injective
 814 $y_H(v) \notin \{y_H(v_1), y_H(v_2)\}$ and belongs to P_y^H , which contradicts the minimality of P_y^H . This
 815 proves Claim 2.

816 Before we describe how to obtain a witness T' of T for b , we must make an observation.
 817 We note that $Im(x_L) \cup Im(x_R) = V(T)$: the idea is that every node of T must originate
 818 from either T_L or T_R .

819 »»»> e150fdde332112fd1c2acb6bd85a9a5606b79547 Now we are able to describe how to
 820 obtain a witness T' of T for b . For every $y \in Y_L^*$, in the last edge of path Q_y^L we plug in the
 821 single pair $(Y_{y'}^L, S_{y'}^L)$ rooted at y' , for every internal node y' of P_y^L , in the order the nodes y'
 822 appear in P_y^L . Note that, in the case an element of Y_L^* is present in more than one P_y^L , we
 823 plug in the corresponding single pair only once. Note also that whenever we plug in some
 824 single pair (Y_y^L, S_y^L) in a DT, the tree Y_y^L has real features and future features as nodes. Call
 825 this graph T^* . Now we do the same sequence of plug ins of the single pairs corresponding to
 826 the internal vertices of P_y^R in the last edge of the path Q_y^R . Again, in the case an element
 827 of Y_R^* is present in more than one P_y^R , we plug in the corresponding single pair only once.
 828 Call the tree obtained in this way T' . Node that T' contains real features from $feat(b_L)$ and
 829 from $feat(b_R)$ and future features with labels in $\mathcal{P}([k])$.

830 To conclude this part of the proof we have to show two things: (i) T is obtained from T'
 831 after removing s vertices; (ii) T' is a real DT for b . We start proving (i): by construction T'
 832 is obtained from T after adding s_L elements from T'_L and s_R elements from T'_R , and so with
 833 $s_L + s_R = s$ more elements.

834 Before considering statement (ii), we consider the following relabelling p_+ of T' : every
 835 real feature in $feat(b_R)$ is assigned to a feature with its label at node b_R and every other
 836 feature is assigned to itself. The real DT T'_L can be obtained from T' by the application of
 837 the composition $r \circ p_* \circ p_+$.

838 Now we consider statement (ii). We show that given an example $e \in exam(b_L)$, e is
 839 correctly classified by T' and to do so we show that e ends in a leaf of T' that corresponds
 840 to the leaf where e ends in T'_L . Say that e goes along a path P of T'_L from the root to a
 841 leaf ℓ and let Q be the corresponding path in T' , i.e. the path from r to ℓ (note that by
 842 construction ℓ is present in T' and is still a leaf). Let v be a node of Q , we can have the
 843 following different cases.

- 844 ■ v is a real feature from $feat(b_L)$: v is also present in T'_L as real feature;
- 845 ■ v is a real feature from $feat(b_R)$: v might not be present in T'_L due reductions but if it is
 present it is a future feature A_i for some $i \in [k]$;
- 847 ■ v is a future feature f_A : v might not be present in T'_L due reductions but if it is present
 it is still the same future feature A_i .

849 If v is present in T'_L then the behaviour of v on e in T'_L and in T' is the same. Suppose
 850 now v is a node of Q that is being reduced due his label and so it is not present in T'_L .
 851 This means there is a set of ancestors of v such that their labels allows to remove v and by
 852 construction v behaves on e like those ancestors. This proves e goes along Q and in particular
 853 it ends at leaf ℓ and so T' is a real DT for b_L . With symmetric construction, we show that
 854 T' is also a real DT for b_R .

855 Now we prove the backward direction. Let T be a reduced DT such that s is the minimum
 856 number of elements that have been deleted from a witness T' of T for b . In particular, we
 857 recall that T' is a real DT for b with actual feature labels in $[k] \cup [k']$ and future feature
 858 labels in $\mathcal{P}([k])$.

859 We create a real DT T'_L by the application of the composition $r \circ p_* \circ p_+$ to T' . By
 860 assumption T' is a real DT for b_L and by construction T'_L is a real DT for b_L . Denote
 861 with T_L the DT template obtained from T'_L by standard reduction and denote with s_L

862 the number of nodes that have been deleted from T'_L to obtain T . By induction we have
 863 $(T_L, s_L) \in \mathcal{R}(b_L)$. Now we note that T_L is obtained from T after the application of the
 864 composition $r \circ p_*$. In a symmetric way, we construct T'_R , T_R and the record $(T_R, s_R) \in \mathcal{R}(b_R)$.
 865 Then $(T, s_L + s_R) \in \mathcal{R}(b)$. \blacktriangleleft

866 ▶ **Lemma 20** (relabel node). *Let $b \in V(B)$ be relabel node. Then $\mathcal{R}(b)$ can be computed in
 867 time $\mathcal{O}(k(2k + 2^k + 2)2^{3k+1})$.*

868 **Proof.** Let b_C be the unique child of b in B . Let R be the mapping of $[k]$ to itself that
 869 represent the node b . Moreover, since we are considering a *nice* NLC-expression we can
 870 assume R is the identity mapping, i.e. $R(\ell) = \ell$, for all values except for a unique element i
 871 of its domain, i.e. $R(i) = j$ for some $j \in [k] \setminus \{i\}$.

872 We say that a future feature A is *good* if it does not distinguish between i and j , that
 873 is $i \in A$ if and only if $j \in A$, and *bad* otherwise. Let (T_C, s_C) be an element of $\mathcal{R}(b_C)$. Let
 874 p'' the following relabelling of the DT template T_C : every feature with label i is assigned
 875 to label j and every future feature with label A is assigned to the future feature with label
 876 $A \setminus \{i\}$.

877 If T_C has a bad future feature then we do not take any other action. Suppose now T_C
 878 has only good future features; now let T be the DT template obtained from T_C after the
 879 application of the composition $r \circ p''$ and let s^* be the number of nodes that have been
 880 deleted from T_C to T .

881 If there is a record in $\mathcal{R}(b)$ of the form (T, s') for some integer $s' \leq s_C + s^*$ then we do
 882 not take any other action. If there is a record in $\mathcal{R}(b)$ of the form (T, s') for some integer
 883 $s' > s_C + s^*$ then we replace it with $(T, s_C + s^*)$. If there is no record in $\mathcal{R}(b)$ of the form
 884 (T, s') for some integer s' then we add $(T, s_C + s^*)$ to $\mathcal{R}(b)$.

885 Now we want to evaluate the running time of computing $\mathcal{R}(b)$. Consider record (T_C, s_C)
 886 in $\mathcal{R}(b_C)$. In $\mathcal{O}(k)$ time we check if T_C all the future features are good. For every such DT
 887 T_C , there are at most 2^{2k} paths from the root to the leaves and for every of these paths there
 888 are at most k nodes for each of the following: feature with label i and and future feature
 889 that contains i . This means $r \circ p''$ can be done in $\mathcal{O}(k)$ time. This means to compute $\mathcal{R}(b)$
 890 takes $\mathcal{O}(k|\mathcal{R}(b_C)|) = \mathcal{O}(k(2k + 2^k + 2)2^{3k+1})$ time.

891 Now we have to show the correctness of the construction for $\mathcal{R}(b)$, i.e. $(T, s) \in \mathcal{R}(b)$ if
 892 and only if s is the minimum number of elements that have been deleted from a witness T'
 893 of T for b .

894 We start with the forward direction. Let $(T, s) \in \mathcal{R}(b)$. By construction there exists a
 895 record $(T_C, s_C) \in \mathcal{R}(b_C)$ such that T is obtained from T_C after the application of $r \circ p''$ and
 896 let $s^* = s - s_C$. By induction s_C is the minimum amount of nodes that have been deleted
 897 from a witness T'_C of T_C for b_C . By construction we also know that every future feature of
 898 both T'_C and T_C is good.

899 Denote with T' the real DT obtained T'_C after the application of $r \circ p''$: note that this
 900 last reduction does not any node since every future feature of T'_C is good and there is no
 901 feature with label i . To conclude this part of the proof we have to show two things: (i) T is
 902 obtained from T' after removing s vertices; (ii) T' is a witness of T for b .

903 Before proving (i), we describe how T can be obtained from T' . Let p''' be the following
 904 relabelling of T' : every real feature that contains j is assigned to the real feature $A \cup \{i\}$
 905 and every other feature is assigned to itself. Then the application of the composition p''' ,
 906 the standard reduction and $r \circ p''$ to T' is exactly the standard reduction for T' which then
 907 result to the DT template T . By Lemma 15 the score of the standard reduction from T' to
 908 T is exactly $s_C + s^* = s$.

909 Now we consider statement (ii). First note that $\text{exam}(b) = \text{exam}(b_C)$. We show that
 910 a given example $e \in \text{exam}(b)$ is correctly classified by T' . Say that e goes along a path P
 911 of T'_C from the root to a leaf ℓ . We show e goes along the path P in T' as well: every real
 912 feature has not changed and so e behaves the same. Since every future feature of T'_C is good,
 913 then e behave the same on the corresponding future feature of T' .

914 Now we prove the backward direction. Let T be a reduced DT such that s is the minimum
 915 number of elements that have been deleted from a witness T' of B for b . In particular, we
 916 recall that real T' is a DT for b with real features and future feature labels in $\mathcal{P}([k] \setminus \{i\})$.

917 We create the real DT T'_C as the application of $r \circ p'''$ to T' , the DT template T_C as the
 918 application of the standard reduction to T'_C . By construction we have $(T_C, s_C) \in \mathcal{R}(b_C)$,
 919 where s_C is the number of nodes that have been removed from T'_C to T_C . Note that T_C has
 920 only good future features. Finally we note that T is obtained from T_C by the application of
 921 $r \circ p''$. \blacktriangleleft

922 Now we can finally prove Theorem 4 and Theorem ??, which we restate here.

923 **Theorem 4 (restated).** *Let E be a CI, let (B, χ) be an NLC-expression decomposition of
 924 width k for $G_I(E)$, and let s be an integer. Then, deciding whether E has a DT of size at
 925 most s is fixed-parameter tractable parameterized by k . In particular, such computation takes
 926 $\mathcal{O}()$ time.*

927 **Proof.** We start off by computing $\mathcal{R}(b)$ for every node b of B , via leaf-to-root dynamic
 928 programming. An upper bound for the running time for this step is the number of nodes of
 929 B times the maximum running time to compute the record at each node which is given by
 930 Lemmas 18, 19 and 20.

931 Now we look at the root node r of B . We go through all the records of $\mathcal{R}(r)$ and select a
 932 record $(T, s) \in \mathcal{R}(r)$ such that $|T| + s$ is minimum over all DTs with no future feature. \blacktriangleleft

933 **Theorem ?? (restated).** DTS is fixed-parameter tractable parameterized by NLC-width.

934 4 An FPT-Algorithm for bounded solution size and δ_{max} .

935 In the following, let E be a CI and $q \notin \text{feat}(E)$. A *pattern* is a pair (T, φ) where $T =$
 936 $(V(T), A(T))$ is a rooted subcubic tree, every leaf-node is either a *positive* or *negative* leaf
 937 and φ maps every inner-node $v \in V(T)$ to an element of $\text{feat}(E) \cup \{q\}$. We say that φ is
 938 the *trait* function of T and $\varphi(v)$ is the trait of node $v \in V(T)$.

939 A pattern (T, φ^*) is an *improvement* for a pattern (T, φ) if $\varphi^*(v) = \varphi(v)$ for every v such
 940 that $\varphi(v) \in \text{feat}(E)$. A *complete improvement* (T, φ^*) for (T, φ) is an improvement such
 941 that $\text{Im}(\varphi^*) \subseteq \text{feat}(E)$. Note that any complete improvement of a pattern is a decision tree.
 942 Given a pattern (T, φ) , an *threshold assignment* of (T, φ) is a function ψ that maps every
 943 node $v \in V(T)$ such that $\varphi(v) \in \text{feat}(E)$ to a rational number $\psi(v)$.

944 Given a threshold assignment ψ for a decision tree (T, φ) , for each node v of T we define
 945 the set of examples that arrives at node v , $E_T(v)$ as follows: $E_T(v)$ is the set of all examples
 946 $e \in E$ such that for each left (right, respectively) arc (u, w) on the unique path from the
 947 root of T to v we have $(\varphi(u))(e) \leq \psi(u)$ ($(\varphi(u))(e) > \psi(u)$, respectively). A decision tree
 948 (T, φ) *correctly classifies* an example $e \in E$ given ψ if e is a positive (negative) example and
 949 $e \in E_T(v)$ for a positive (negative) leaf. We say that (T, φ) *classifies* E given ψ if T correctly
 950 classifies every example $e \in E$ given ψ .

951 We say that (T, φ) can *classify* E if there exists a complete improvement (T, φ^*) for (T, φ)
 952 and there exists a threshold assignment ψ for (T, φ^*) such that (T, φ^*) classifies E given ψ .

953 **4.1 Preprocess**

954 Let E be a CI, and (T, φ) be a pattern. For every $v \in V(T)$, we define the set of *expected examples* E_v as follows:

- 956 ■ if v is the root, then $E_v = E$;
- 957 ■ if v is the left child of a node v_p such that $\varphi(v_p) \in \text{feat}(E)$, then $E_v = E_{v_p}[\varphi(v_p) \leq th_L(v_p) + 1]$;
- 959 ■ if v is the right child of a node v_p such that $\varphi(v_p) \in \text{feat}(E)$, then $E_v = E_{v_p}[\varphi(v_p) > th_R(v_p) - 1]$;
- 961 ■ if v is a child of a v_p such that $\varphi(v_p) = q$, then $E_v = E_{v_p}$.

962 Note that the definition of E_v is strictly related with the following: if $\varphi(v) \in \text{feat}(E)$,
963 let c_ℓ and c_r be the left, resp. right, child of v , we define two values $th_L(v)$ and $th_R(v)$ as
964 follows:

- 965 ■ let $th_L(v)$ be the maximum value in $D_E(\varphi(v))$ such that (T_{c_ℓ}, φ) can classify every example in $E_v[\varphi(v) \leq th_L(v)]$;
- 967 ■ let $th_R(v)$ be the minimum value in $D_E(\varphi(v))$ such that (T_{c_r}, φ) can classify every example in $E_v[\varphi(v) > th_R(v)]$.

969 Before formally proving in Lemma 21 that we are able to compute E_v and $th_L(v)$, $th_R(v)$
970 (when $\varphi(v) \in \text{feat}(E)$) for every $v \in V(T)$, we want to describe the role of E_v in the proof
971 of Lemma 22.

972 Let us consider the following situation. Suppose we are trying to find a DT of minimum size for a CI E using at least the features in a given support set S . The first step would be to compute a minimum size DT T^* for E such that $\text{feat}(T^*) = S$. Next we analyse the case an optimal DT for E uses not only every feature from S but some additional feature: for this reason we consider patterns (T, φ) with T of size at most s and such that $S \cup \{q\} = \text{Im}(\varphi)$.

977 Let us recall a definition. Let (T, φ) be a pattern and $v \in V(T)$ be an inner node of T with left child ℓ , right child r , and parent p . We say that T' is obtained from T after *left/right-contracting* v if T' is a rooted subcubic tree obtained from T after removing v together with all nodes in T_r/T_ℓ and adding the edge between p and l/r ; if v has no parent then no edge is added.

982 In order to argue that a pattern (T, φ) can classify E , we have first to compute a complete improvement (or a series of improvements that ends up in a complete improvement) of (T, φ) . Let $v \in V(T)$ such that $\varphi(v) = q$. If the pattern (T', φ) , where T' is obtained after left/right-contracting v , can classify every example in E_v , then we discard the pattern (T, φ) : (T', φ) is equivalent to (T, φ) but of smaller size. TO EXPAND.

987 **Claim 1.** *Let (T, φ) be a pattern. For every node $v \in V(T)$, (T_v, φ) can not classify every example in E_v .*

989 We prove this claim by induction on the height of T .

990 Let us first consider the case in which v is the root of T . Here we have that $T_v = T$ and $E_v = E$ and so the claim directly follows from the assumption that T can not be extended to a DT for E . Now let v be the child of an unknown feature node v_p . Recall that, by assumption, v_p does not distinguish any pair of examples and so it is safe to assume that E_{v_p} is a set of examples that can not be classified by extending T_v .

995 Finally suppose now that v is the left child of a supported feature node v_p . By the maximality of $th_L(v_p)$, the set $E_v = E_{v_p}[f_p \leq th_L(v_p) + 1]$ necessarily contains examples that can not be classified by extending T_v . In a similar way, the minimality of $th_R(v)$ shows

998 that the claim holds when v is the right child of a supported feature node. This proves the
 999 Claim 1.

1000 Thanks to Claim 1, it follows that for every node v that has a supported feature, $th_L(v) <$
 1001 $th_R(v)$. Suppose that for some node v with supported feature f we have $th_L(v) \geq th_R(v)$.
 1002 Setting $(feat(v) = th_L(v))$ generates a DT for E_v , which is a contradiction.

1003 ▶ **Lemma 21.** *Let E be a CI, let T be a partial DT of depth at most d and let $\bar{\alpha}$ be a
 1004 S -feature assignment of T for a support set S . Then there is an algorithm that runs in time
 1005 $\mathcal{O}(2^{d^2/2} n^{1+o(1)} \log n)$ and computes the set E_v and thresholds $th_L(v)$ and $th_R(v)$ for every
 1006 node $v \in V(T)$.*

1007 **Proof.** The idea is to use the recursive algorithm **findLR** illustrated in Algorithm 1. That
 1008 is, given E , T , $\bar{\alpha}$, the algorithm **findLR** attempts to find the triples $(E_v, th_L(v), th_R(v))$
 1009 for every node $v \in V(T)$. Lines 3 to 4: if T consists of a leaf node, the algorithm just
 1010 report $(E, \text{nil}, \text{nil})$. Let c_ℓ and c_r be the left, resp. right, child of the root v . Lines 6
 1011 to 11: if the root of T has an unknown feature, the algorithm calls itself recursively to
 1012 compute the triple for (E, T_{c_ℓ}, α) and (E, T_{c_r}, α) . Lines 13 to 15: if the root of T has a
 1013 supported feature f , the algorithm computes the pair (t_ℓ, t_r) for the root using the algorithm
 1014 **binarySearch** and then calls itself recursively to compute the triple for $(E[f \leq t_\ell + 1], T_{c_\ell}, \alpha)$
 1015 and $(E[f > t_r - 1], T_{c_r}, \alpha)$.

1016 A key element for the correctness of **findLR** is the algorithm **binarySearch** illustrated
 1017 in Algorithm 2. Given E , T , α , f , c_ℓ and c_r , this algorithm computes the pair (t_ℓ, t_r) for
 1018 the root of T . This sub-routine performs a standard binary search procedure on the array
 1019 D containing all the values in $D_E(f)$ in ascending order to find maximum t_ℓ and minimum
 1020 t_r such that (T_{c_ℓ}, α) and (T_{c_r}, α) can be extended to DT for $E[f \leq t_\ell]$ and for $E[f > t_r]$
 1021 respectively. To achieve this, the sub-routine makes at most $\log |E|$ calls to **findTH**; note
 1022 that each of those calls is made for a tree of smaller depth. Lines 3 to 12: the algorithm
 1023 finds the maximum t_ℓ by calling algorithm **findTH** in Line 6 repeatedly. Lines 13 to 22: the
 1024 algorithm finds the minimum t_r by calling algorithm **findTH** in Line 16 repeatedly.

1025 A sub-routine used for **binarySearch** is the algorithm **findTH** illustrated in Algorithm 3.
 1026 This algorithm is very similar to Algorithm 1 but the output is some way much simpler.

1027 The running time of Algorithm 1 can now be obtained by multiplying the number of
 1028 recursive calls to **findLR** with the time required for one recursive call. To obtain the number
 1029 of recursive calls first note that if **findLR** is called with a pseudo DT of depth d , then it
 1030 makes at most $(2 \log n) + 2$ recursive calls to **findLR** with a pseudo DT of depth at most
 1031 $d - 1$, where $n = |E|$. Therefore the number $T(n, d)$ of recursive calls for a pseudo DT of
 1032 depth d is given by the recursion relation $T(n, d) = (2(\log n) + 2)T(n, d - 1)$ starting with
 1033 $T(n, 0) = 0$. This implies that $T(n, d) \in \mathcal{O}((\log n)^d)$. Finally, the runtime for one recursive
 1034 call is easily seen to be at most $\mathcal{O}(n \log n)$. Hence, the total runtime of the algorithm is at
 1035 most $\mathcal{O}((\log n)^d n \log n)$, which because (see also [9, Exercise 3.18]):

$$1036 (\log n)^d \leq 2^{d^2/2} 2^{\log \log d^2/2} = 2^{d^2/2} n^{o(1)}$$

1037 is at most $\mathcal{O}(2^{d^2/2} n^{1+o(1)} \log n)$. ◀

1038 4.2 The algorithm

1039 Now we have computed a set E_v for every node $v \in V(T)$, whether it is a leaf or has a
 1040 supported or unknown feature. A *pool set* for node $v \in V(T)$ is a set $\Pi(v) \subseteq E_v$, such that if
 1041 every example of $\Pi(v)$ arrives at node v then either

Algorithm 1 Algorithm to compute the triple $(E_v, th_L(v), th_R(v))$ for every node $v \in V(T)$.

Input: CI E , pseudo DT T , S -feature assignment α for T
Output: a triple $(E_v, th_L(v), th_R(v))$ for every node $v \in V(T)$.

```

1: function FINDLR( $E, T, \alpha$ )
2:    $r \leftarrow$  “root of  $T$ ”
3:   if  $r$  is a leaf then
4:     return  $(E, \text{nil}, \text{nil})$ 
5:    $c_\ell, c_r \leftarrow$  “left child and right child of  $r$ ”
6:   if  $r$  has an unknown feature then
7:      $\lambda_\ell \leftarrow$  FINDLR( $E, T_{c_\ell}, \alpha$ )
8:      $\lambda_r \leftarrow$  FINDLR( $E, T_{c_r}, \alpha$ )
9:     if  $\lambda_\ell \neq \text{nil}$  and  $\lambda_r \neq \text{nil}$  then
10:      return  $(E, \text{nil}, \text{nil}) \cup \lambda_\ell \cup \lambda_r$ 
11:    return  $\text{nil}$ 
12:     $f \leftarrow \alpha(r)$ 
13:     $(t_\ell, t_r) \leftarrow$  BINARYSEARCH( $E, T, \alpha, c_\ell, c_r, f$ )
14:     $\lambda_\ell \leftarrow$  FINDLR( $E[f \leq t_\ell + 1], T_{c_\ell}, \alpha$ )
15:     $\lambda_r \leftarrow$  FINDLR( $E[f > t_r - 1], T_{c_r}, \alpha$ )
16:    return  $(E, t_\ell, t_r) \cup \lambda_\ell \cup \lambda_r$ 
```

- 1042 ■ T can not be extended to a DT for E , or
1043 ■ there is an node $u \in V(T_v)$ with unknown feature f_u and two elements $e, e' \in \Pi(v)$ such
1044 that f_u must distinguish e and e' .

1045 For every node $v \in V(T)$, we define $\Pi(v)$ in a leaves-to-root fashion as follows:

- 1046 (a) if v is a negative leaf then $\Pi(v) = \{e^+\}$, where e^+ is any example in $E^+ \cap E_v$;
1047 (b) if v is a positive leaf then $\Pi(v) = \{e^-\}$, where e^- is any example in $E^- \cap E_v$;
1048 (c) if v is a supported feature node. Let c_ℓ and c_r be the left, resp. right, child of v , then
1049 $\Pi(v) = \Pi(c_\ell) \cup \Pi(c_r)$;
1050 (d) if t is an unknown feature node. Let c_ℓ and c_r be the left, resp. right, child of v , then
1051 $\Pi(v) = \Pi(c_\ell) \cup \Pi(c_r)$.

1052 Now we want to show that the construction of Π is correct, that is:

1053 **Claim 2.** $\Pi(v)$ is a pool set for v for every node $v \in V(T)$.

1054 We start proving (a): let v be a negative leaf and suppose there is an element $e^+ \in E^+ \cap E_v$,
1055 then there is no threshold assignment for T that would correctly classify e . The correctness
1056 for (b) is obtained by a symmetric argument. Let us consider (c) and let $f \in S$ be the
1057 feature at v : thanks to the preprocessing, for every $e_\ell \in \Pi(c_\ell)$ and $e_r \in \Pi(c_r)$, we know that
1058 $f(e_\ell) < f(e_r)$. This means that either every element of $\Pi(c_\ell)$ is sent to c_ℓ or every element of
1059 $\Pi(c_r)$ is sent to c_r : by induction the statement is proven. Finally we consider statement (d).
1060 If every element of $\Pi(c_\ell)$ is sent to c_ℓ or every element of $\Pi(c_r)$ is sent to c_r , the statement
1061 is proven by induction. Otherwise, there is an example $e_\ell \in \Pi(c_\ell)$ that ends in c_r and an
1062 example $e_r \in \Pi(c_r)$ that ends in c_ℓ . This means v is an unknown feature node in T_t that
1063 distinguishes two examples in $\Pi(v)$, namely e_ℓ and e_r . This proves Claim 2.

1064 In particular, let us consider the pool set $\Pi(r)$ for the root r of T , we define $\Pi(T) := \Pi(r)$.
1065 In this way given T , we are able to compute the corresponding pool set.

1066 Let S be a support set for a CI E , we stay that $B \subseteq \text{feat}(E)$ is a *branching set* for S if
1067 for every minimal DT T for E such that $S \subset \text{feat}(T)$ then $B \cap (\text{feat}(T) \setminus S) \neq \emptyset$.

■ **Algorithm 2** Algorithm to compute the pair $(th_L(r), th_R(r))$ for the root r of T

Input: CI E , pseudo DT T , S -feature assignment α for T , feature f of the root of T , left child c_ℓ of the root of T , right child c_r of the root of T

Output: maximum threshold t_ℓ in $D_E(f)$ for f such that (T_{c_ℓ}, α) can be extended to a DT for $E[f \leq t_\ell]$ and minimum threshold t_r in $D_E(f)$ for f such that (T_{c_r}, α) can be extended to a DT for $E[f > t_r]$

```

1: function binarySearch( $E, T, \alpha, f, c_\ell, c_r$ )
2:    $D \leftarrow$  “array containing all elements in  $D_E(f)$  in
      ascending order”
3:    $L \leftarrow 0; R \leftarrow |D_E(f)| - 1; b \leftarrow 0$ 
4:   while  $L \leq R$  do
5:      $m \leftarrow \lfloor (L + R)/2 \rfloor$ 
6:     if FINDTH( $E[f \leq D[m]], T_{c_\ell}, \alpha$ ) = TRUE then
7:        $L \leftarrow m + 1; b \leftarrow 1$ 
8:     else
9:        $R \leftarrow m - 1; b \leftarrow 0$ 
10:    if  $b = 1$  then
11:       $t_\ell \leftarrow D[m]$ 
12:     $t_\ell \leftarrow D[m - 1]$                                 ▷ assuming that  $D[-1] = D[0] - 1$ 
13:     $L \leftarrow 0; R \leftarrow |D_E(f)| - 1; b \leftarrow 0$ 
14:    while  $L \leq R$  do
15:       $m \leftarrow \lfloor (L + R)/2 \rfloor$ 
16:      if FINDTH( $E[f > D[m]], T_{c_r}, \alpha$ ) = TRUE then
17:         $R \leftarrow m - 1; b \leftarrow 1$ 
18:      else
19:         $L \leftarrow m + 1; b \leftarrow 0$ 
20:    if  $b = 1$  then
21:       $t_r \leftarrow D[m]$ 
22:     $t_r \leftarrow D[m + 1]$                                 ▷ assuming that  $D[|D_E(f)|] = D[|D_E(f)| - 1] + 1$ 
23:  return  $(t_\ell, t_r)$ 

```

1068 ► **Lemma 22.** There is a $\mathcal{O}(2^{d^2/2} s^{2s+1} n^{1+o(1)} \log n)$ time algorithm that given a support set
 1069 S computes a branching set R_0 for S of size at most $s^{2s+3}\delta_{\max}$.

1070 **Proof.** Let E be a CI, a support set S for E and an integer s . We start by enumerating all
 1071 partial DTs of size at most s such that the set of supported features is exactly S . For every
 1072 such partial DT T , thanks to Lemma 21, we are able to obtain the set E_v for every node
 1073 $v \in V(T)$ in time $\mathcal{O}(2^{d^2/2} n^{1+o(1)} \log n)$. In a leaves-to-root fashion, we are able to compute
 1074 the set $\Pi(v)$ for every node $v \in V(T)$ and ultimately $\Pi(T)$.

1075 Let $R(T)$ be the set of all the features in $feat(E) \setminus S$ that distinguish at least two
 1076 examples in $\Pi(T)$. The algorithm returns the set of features R_0 obtained by considering the
 1077 union of the sets $R(T)$ over all these partial DTs T for S of size at most s . By Lemma 1 this
 1078 algorithm runs in time $\mathcal{O}(2^{d^2/2} s^{2s+1} n^{1+o(1)} \log n)$.

1079 Now we show the size of R_0 is bounded. By construction $|\Pi(T)| \leq |T| \leq s$; for every two
 1080 distinct elements of $\Pi(T)$, by definition, there are at most δ_{\max} features that distinguish
 1081 such two examples. This means that $|R(T)| \leq s^2\delta_{\max}$ and so R_0 has size at most $s^{2s+3}\delta_{\max}$.

1082 We are left to show that R_0 is a branching set for S . Let T be a minimal DT for E such
 1083 that $S \subset feat(T)$ and suppose by contradiction that $R_0 \cap (feat(T) \setminus S) = \emptyset$. In particular we
 1084 have that $R(T) \cap (feat(T) \setminus S) = \emptyset$. This means that every unknown feature of T does not
 1085 distinguish any two elements in $\Pi(T)$. By Claim 2, $\Pi(T) = \Pi(r)$, where r is the root of T , is
 1086 a pool set and so T can not be extended to a DT for E , which is a contradiction. ◀

■ **Algorithm 3** Algorithm to compute the threshold assignment for a pseudo DT and a feature assignment.

Input: CI E , pseudo DT T , S -feature assignment α for T
Output: TRUE if all examples in E can be correctly classified by (T, α) , FALSE otherwise

```

1: function findTH( $E, T, \alpha$ )
2:    $r \leftarrow$  “root of  $T$ ”
3:   if  $r$  is a leaf then
4:     if  $E$  is not uniform then
5:       return FALSE
6:     return TRUE
7:    $c_\ell, c_r \leftarrow$  “left child and right child of  $r$ ”
8:   if  $r$  has an unknown feature then
9:      $\lambda_\ell \leftarrow \text{FINDTH}(E, T_{c_\ell}, \alpha)$ 
10:     $\lambda_r \leftarrow \text{FINDTH}(E, T_{c_r}, \alpha)$ 
11:    if  $\lambda_\ell = \text{TRUE}$  and  $\lambda_r = \text{TRUE}$  then
12:      return TRUE
13:    return FALSE
14:    $f \leftarrow \alpha(r)$ 
15:    $t \leftarrow \text{BINARYSEARCH}(E, T, \alpha, c_\ell, f)$ 
16:    $\lambda_r \leftarrow \text{FINDTH}(E[f > t], T_{c_r}, \alpha)$ 
17:   if  $\lambda_r = \text{FALSE}$  then
18:     return FALSE
19:   return TRUE

```

1087 ▶ **Lemma 23** ([23]). Let E be a CI and let k be an integer. Then there is an algorithm that
1088 in time $\mathcal{O}(\delta_{\max}(E)^k |E|)$ enumerates all (of the at most $\delta_{\max}(E)^k$) minimal support sets of
1089 size at most k for E .

1090 ▶ **Lemma 24** ([23]). Let T be a DT of minimum size for E and let S be a support set
1091 contained in $\text{feat}(T)$. Then, the set $R = \text{feat}(T) \setminus S$ is useful.

1092 ▶ **Theorem 25.** MINIMUM DECISION TREE SIZE is fixed-parameter tractable parametrized
1093 by $\delta_{\max} + s$.

1094 **Proof.** We start by presenting the algorithm for MINIMUM DECISION TREE SIZE, which is
1095 illustrated in Algorithm 4 and Algorithm 5.

1096 Given a CI E and an integer s , the algorithm returns a DT of minimum size among all
1097 DTs of size at most s if such a DT exists and otherwise the algorithm returns **nil**. The
1098 algorithm **minDT** starts by computing the set \mathcal{S} of all minimal support sets for E of size
1099 at most s , which because of Lemma 23 results in a set \mathcal{S} of size at most $\binom{|E|}{s}$. In Line 4
1100 the algorithm then iterates over all sets S in \mathcal{S} and calls the function **minDTS** given in
1101 Algorithm 5 for E , s , and S , which returns a DT of minimum size among all DTs T for E
1102 of size at most s such that $S \subseteq \text{feat}(T)$. It then updates the currently best decision tree B
1103 if necessary with the DT found by the function **minDTS**. Moreover, if the best DT found
1104 after going through all sets in \mathcal{S} has size at most s , it is returned (in Line 9), otherwise
1105 the algorithm returns **nil**. Finally, the function **minDTS** given in Algorithm 5 does the
1106 following. It first computes a DT T of minimum size that uses exactly the features in S
1107 using Lemma ???. It then tries to improve upon T with the help of useful sets. That is, it
1108 uses Lemma 22 to compute the branching set R_0 . It then iterates over all (of the at most
1109 $\binom{|S|}{1}$) features $f \in R_0$ (using the for-loop in Line 4), and calls itself recursively on the feature
1110 set $S \cup \{f\}$. If this call finds a smaller DT, then the current best DT B is updated. Finally,
1111 after the for-loop the algorithm either returns B if its size is less than s or **nil** otherwise.

1112 Towards showing the correctness of Algorithm 4, consider the case that E has a DT
 1113 of size at most s and let T be a such a DT of minimum size. Because of Observation ??,
 1114 $\text{feat}(T)$ is a support set for E and therefore $\text{feat}(T)$ contains a minimal support set S of size
 1115 at most s . Because the for-loop in Line 4 of Algorithm 4 interates over all minimal support
 1116 sets of size at most s for E , it follows that Algorithm 5 is called with parameters E , s , and
 1117 S . If $\text{feat}(T) = S$, then B is set to a DT for E of size $|T|$ in Line 2 of Algorithm 5 and the
 1118 algorithm will output a DT of size at most $|T|$ for E . If, on the other hand, $\text{feat}(T) \setminus S \neq \emptyset$,
 1119 then because T has minimum size and S is a support set for E with $S \subseteq \text{feat}(T)$, we obtain
 1120 from Lemma 24 that the set $R = \text{feat}(T) \setminus S$ is useful for S . Therefore, because of Lemma 22,
 1121 R has to contain a feature f from the set R_0 computed in Line 3. It follows that Algorithm 5
 1122 is called with parameters E , s , and $S \cup \{v\}$. From now onwards the argument repeats and
 1123 since $R_0 \neq \emptyset$ the process stops after at most $s - |S|$ recursive calls after which a DT for E of
 1124 size at most $|T|$ will be computed in Line 2 of Algorithm 5. Finally, it is easy to see that if
 1125 Algorithm 4 outputs a DT T , then it is a valid solution. This is because, T must have been
 1126 computed in Line 2 of Algorithm 5, which implies that T is a DT for E . Moreover, T has
 1127 size at most s , because of Line 8 in Algorithm 4.

1128 To analyse the run-time of the algorithm, we first remark that the whole algorithm can
 1129 be seen as a bounded-depth search tree algorithm, i.e., a branching algorithm with small
 1130 recursion depth and few branches at every node. In particular, every recursive call adds at
 1131 least one feature to the set of features bounding the recursion depth to at most s . Moreover,
 1132 every feature that is added is either added in Line 2 of Algorithm 4, when enumerating
 1133 all minimal support sets, in which case there are at most $\delta_{\max}(E)$ branches or the feature
 1134 is added in Line 5 of Algorithm 5, in which case there are at most $|R_0| \leq s^{2s+3}\delta_{\max}(E)$
 1135 branches. It follows that the algorithm can be seen as a branching algorithm of depth
 1136 at most s with at most $s^{2s+3}\delta_{\max}(E) = \max\{s^{2s+3}\delta_{\max}(E), \delta_{\max}(E)\}$ branches at every
 1137 step. Therefore, the total run-time of the algorithm is at most the number of nodes in
 1138 the branching tree, i.e., at most $(s^{2s+3}\delta_{\max}(E))^s$, times the maximum time required in
 1139 one recursive call. Now the maximum time required for one recursive call is dominated
 1140 by the time spend in Line 2 of Algorithm 5, i.e., the time required to compute a DT of
 1141 minimum size using exactly the features in S with the help of Theorem ??, which is at
 1142 most $2^{\mathcal{O}(s^2)}\|E\|^{1+o(1)}\log\|E\|$. Therefore, we obtain $(s^{2s+3}\delta_{\max}(E))^s 2^{\mathcal{O}(s^2)}\|E\|^{1+o(1)}\log\|E\|$
 1143 as the total run-time of the algorithm, which shows that DTS is fixed-parameter tractable
 1144 parameterized by $s + \delta_{\max}(E)$. \blacktriangleleft

Algorithm 4 Main method for finding a DT of minimum size.

Input: CI E and integer s
Output: DT for E of minimum size (among all DTs of size at most s) if such a DT exists, otherwise
 nil

```

1: function minDT( $E$ ,  $s$ )
2:    $S \leftarrow$  "set of all minimal support sets for  $E$  of size at most  $s$  using Lemma 23"
3:    $B \leftarrow$  nil
4:   for  $S \in S$  do
5:      $T \leftarrow$  MINDTS( $E$ ,  $s$ ,  $S$ )
6:     if ( $T \neq$  nil) and ( $B =$  nil or  $|B| > |T|$ ) then
7:        $B \leftarrow T$ 
8:     if  $B \neq$  nil and  $|B| \leq s$  then
9:       return  $B$ 
10:  return nil
  
```

Algorithm 5 Method for finding a DT of minimum size using at least the features in a given support set S .

Input: CI E , integer s , support set S for E with $|S| \leq s$
Output: DT of minimum size among all DTs T for E of size at most s such that $S \subseteq \text{feat}(T)$; if no such DT exists, **nil**

```

1: function minDTS( $E, s, S$ )
2:    $B \leftarrow$  “compute a DT of minimum size for  $E$  using exactly the features in  $S$  using Theorem ??”
3:    $R_0 \leftarrow$  “compute the branching set  $R_0$  for  $S$  using Lemma 22”
4:   for  $f \in R_0$  do
5:      $T \leftarrow \text{MINDTS}(E, s, S \cup \{f\})$ 
6:     if  $T \neq \text{nil}$  and  $|T| < |B|$  then
7:        $B \leftarrow T$ 
8:     if  $|B| \leq s$  then
9:       return  $B$ 
10:    return nil

```

1145 **5 Conclusion**

1146 We have initiated the study of the parameterized complexity of learning DTs from data. Our
1147 main tractability result provides novel insights into the structure of DTs and is based on
1148 the NLC-width parameter that seems to be well suited to measure the complexity of input
1149 instances for the problem.

1150 The problem of learning DTs comes in many variants and flavors, which opens up a wide
1151 range of new research directions to explore. For instance:

- 1152 ■ What other (structural) parameters can be exploited to efficiently learn DTs? Is learning
1153 DTs of small size fixed-parameter tractable parameterized by the rank-width of $G_I(E)$?
- 1154 ■ Instead of learning DTs of small size, one often wants to learn DTs of small height.
1155 Therefore, it is natural to ask whether our approach can be also used in this setting.
1156 While one can adapt our approach to obtain an XP-algorithm for learning DTs of small
1157 height parameterized by NLC-width, it is not clear to us whether the problem also allows
1158 for an fpt-algorithm.
- 1159 ■ Can we extend our approach to CIs, where features range over an arbitrary domain? In
1160 this case, one usually still uses DTs that make binary decisions (i.e. whether a feature is
1161 smaller equal or larger than a given threshold). While it is relatively easy to see that our
1162 approach can be extended if the domain’s size (for every feature) is bounded or used as
1163 an additional parameter, it is not clear what happens if the size of the domain is allowed
1164 to grow arbitrarily.

1165 **References**

- 1166 1 Jørgen Bang-Jensen and Gregory Gutin. *Digraphs*. Springer Monographs in Mathematics.
1167 Springer-Verlag London Ltd., London, second edition, 2009.
- 1168 2 Christian Bessiere, Emmanuel Hebrard, and Barry O’Sullivan. Minimising decision tree size
1169 as combinatorial optimisation. In Ian P. Gent, editor, *Principles and Practice of Constraint
1170 Programming - CP 2009*, pages 173–187, Berlin, Heidelberg, 2009. Springer Berlin Heidelberg.
- 1171 3 Hans L. Bodlaender. A linear-time algorithm for finding tree-decompositions of small treewidth.
1172 *SIAM J. Comput.*, 25(6):1305–1317, 1996.
- 1173 4 Endre Boros, Yves Crama, Peter L. Hammer, Toshihide Ibaraki, Alexander Kogan, and
1174 Kazuhisa Makino. Logical analysis of data: classification with justification. *Ann. Oper. Res.*,
1175 188(1):33–61, 2011.

- 1176 5 Endre Boros, Vladimir Gurvich, Peter L. Hammer, Toshihide Ibaraki, and Alexander Kogan.
1177 Decomposability of partially defined Boolean functions. *Discr. Appl. Math.*, 62(1-3):51–75,
1178 1995.
- 1179 6 Endre Boros, Takashi Horiyama, Toshihide Ibaraki, Kazuhisa Makino, and Mutsunori Yagiura.
1180 Finding essential attributes from binary data. *Ann. Math. Artif. Intell.*, 39:223–257, 11 2003.
1181 doi:10.1023/A:1024653703689.
- 1182 7 Endre Boros, Toshihide Ibaraki, and Kazuhisa Makino. Variations on extending partially
1183 defined Boolean functions with missing bits. *Information and Computation*, 180(1):53–70,
1184 2003.
- 1185 8 Yves Crama, Peter L. Hammer, and Toshihide Ibaraki. Cause-effect relationships and partially
1186 defined Boolean function. *Ann. Oper. Res.*, 16:299–326, 1988.
- 1187 9 Marek Cygan, Fedor V. Fomin, Lukasz Kowalik, Daniel Lokshtanov, Dániel Marx, Marcin
1188 Pilipczuk, Michal Pilipczuk, and Saket Saurabh. *Parameterized Algorithms*. Springer, 2015.
1189 doi:10.1007/978-3-319-21275-3.
- 1190 10 Adnan Darwiche and Auguste Hirth. On the reasons behind decisions. In Giuseppe De
1191 Giacomo, Alejandro Catalá, Bistra Dilkina, Michela Milano, Senén Barro, Alberto Bugarín,
1192 and Jérôme Lang, editors, *ECAI 2020 - 24th European Conference on Artificial Intelligence,
1193 29 August-8 September 2020, Santiago de Compostela, Spain, August 29 - September 8, 2020 -
1194 Including 10th Conference on Prestigious Applications of Artificial Intelligence (PAIS 2020)*,
1195 volume 325 of *Frontiers in Artificial Intelligence and Applications*. IOS Press, 2020.
- 1196 11 Reinhard Diestel. *Graph Theory*, volume 173 of *Graduate Texts in Mathematics*. Springer
1197 Verlag, New York, 2nd edition, 2000.
- 1198 12 Finale Doshi-Velez and Been Kim. A roadmap for a rigorous science of interpretability. *CoRR*,
1199 abs/1702.08608, 2017. URL: <http://arxiv.org/abs/1702.08608>, arXiv:1702.08608.
- 1200 13 Rodney G. Downey and Michael R. Fellows. *Fundamentals of parameterized complexity*. Texts
1201 in Computer Science. Springer Verlag, 2013.
- 1202 14 Wolfgang Espelage, Frank Gurski, and Egon Wanke. Deciding clique-width for graphs of
1203 bounded tree-width. *J. Graph Algorithms Appl.*, 7(2):141–180, 2003.
- 1204 15 Bryce Goodman and Seth R. Flaxman. European union regulations on algorithmic decision-
1205 making and a “right to explanation”. *AI Magazine*, 38(3):50–57, 2017.
- 1206 16 Laurent Hyafil and Ronald L. Rivest. Constructing optimal binary decision trees is NP-complete.
1207 *Information Processing Letters*, 5(1):15–17, 1976.
- 1208 17 Toshihide Ibaraki, Yves Crama, and Peter L. Hammer. *Partially defined Boolean functions*,
1209 page 511–563. Encyclopedia of Mathematics and its Applications. Cambridge University Press,
1210 2011.
- 1211 18 Daniel T. Larose. *Discovering knowledge in data*. Wiley-Interscience [John Wiley & Sons],
1212 Hoboken, NJ, 2005. An introduction to data mining.
- 1213 19 Zachary C. Lipton. The mythos of model interpretability. *Communications of the ACM*,
1214 61(10):36–43, 2018.
- 1215 20 E.J. McCluskey. *Introduction to the Theory of Switching Circuits*. Electrical and electronic
1216 engineering series. Princeton University series. McGraw-Hill, 1965.
- 1217 21 Don Monroe. AI, explain yourself. *AI Communications*, 61(11):11–13, 2018. doi:10.1145/
1218 3276742.
- 1219 22 Seerama K. Murthy. Automatic construction of decision trees from data: A multi-disciplinary
1220 survey. *Data Min. Knowl. Discov.*, 2(4):345–389, 1998. doi:10.1023/A:1009744630224.
- 1221 23 Sebastian Ordyniak and Stefan Szeider. Parameterized complexity of small decision tree learn-
1222 ing. In *Proceedings of AAAI'21, the Thirty-Fifth AAAI Conference on Artificial Intelligence*,
1223 pages 6454–6462. AAAI Press, 2021.
- 1224 24 Sang-il Oum. Approximating rank-width and clique-width quickly. *ACM Transactions on
1225 Algorithms*, 5(1), 2008.
- 1226 25 J. Ross Quinlan. Induction of decision trees. *Machine Learning*, 1(1):81–106, 1986. doi:
1227 10.1023/A:1022643204877.

- 1228 **26** Felix Reidl, Peter Rossmanith, Fernando Sánchez Villaamil, and Somnath Sikdar. A faster
1229 parameterized algorithm for treedepth. In Javier Esparza, Pierre Fraigniaud, Thore Husfeldt,
1230 and Elias Koutsoupias, editors, *Automata, Languages, and Programming - 41st International
1231 Colloquium, ICALP 2014, Copenhagen, Denmark, July 8-11, 2014, Proceedings, Part I*, volume
1232 8572 of *Lecture Notes in Computer Science*, pages 931–942. Springer, 2014.
1233 **27** Richard Stanley and Eric W. Weisstein. Catalan number, from mathworld—a wolfram web
1234 resource, 2015.
1235 **28** Egon Wanke. k -NLC graphs and polynomial algorithms. *Discr. Appl. Math.*, 54(2-3):251–266,
1236 1994. Efficient algorithms and partial k -trees.